

# Verilog Implementation of Decision Tree Accelerator

Bo-Fan, Chen

Side Project for 2022 Job Interviews

# Outline

- Problem/ Dataset
- Software
- Conversion
- Hardware

# Problem Description

- Predict payment failure
  - Decision tree
- Decision tree can be used in various fields
  - Some of them require ***high processing speed***
    - One solution is to adopt a hardware accelerator dedicated for decision tree

# Dataset

- Label
  - Payment success/fail, 0/1
- Features
  - 13 features for each application data
  - Including FICO score, interest rate...
- 9578 loaning data
  - 8045 successes / 1533 fails

# Data Pre-process

- “Loaning Purpose” is a categorical feature
  - 6 categories
  - One hot encoding
  - 18 total features
- Normalize the features to 0~255, integer
  - **Decision thresholds will also be 0~255, integer!**
  - Hardware-friendly modification
  - At the cost of potential performance drop
- Training/Testing split
  - Randomly sample
  - 6704/2874

# Outline

- Problem/ Dataset
- **Software**
- Conversion
- Hardware

# Decision Tree

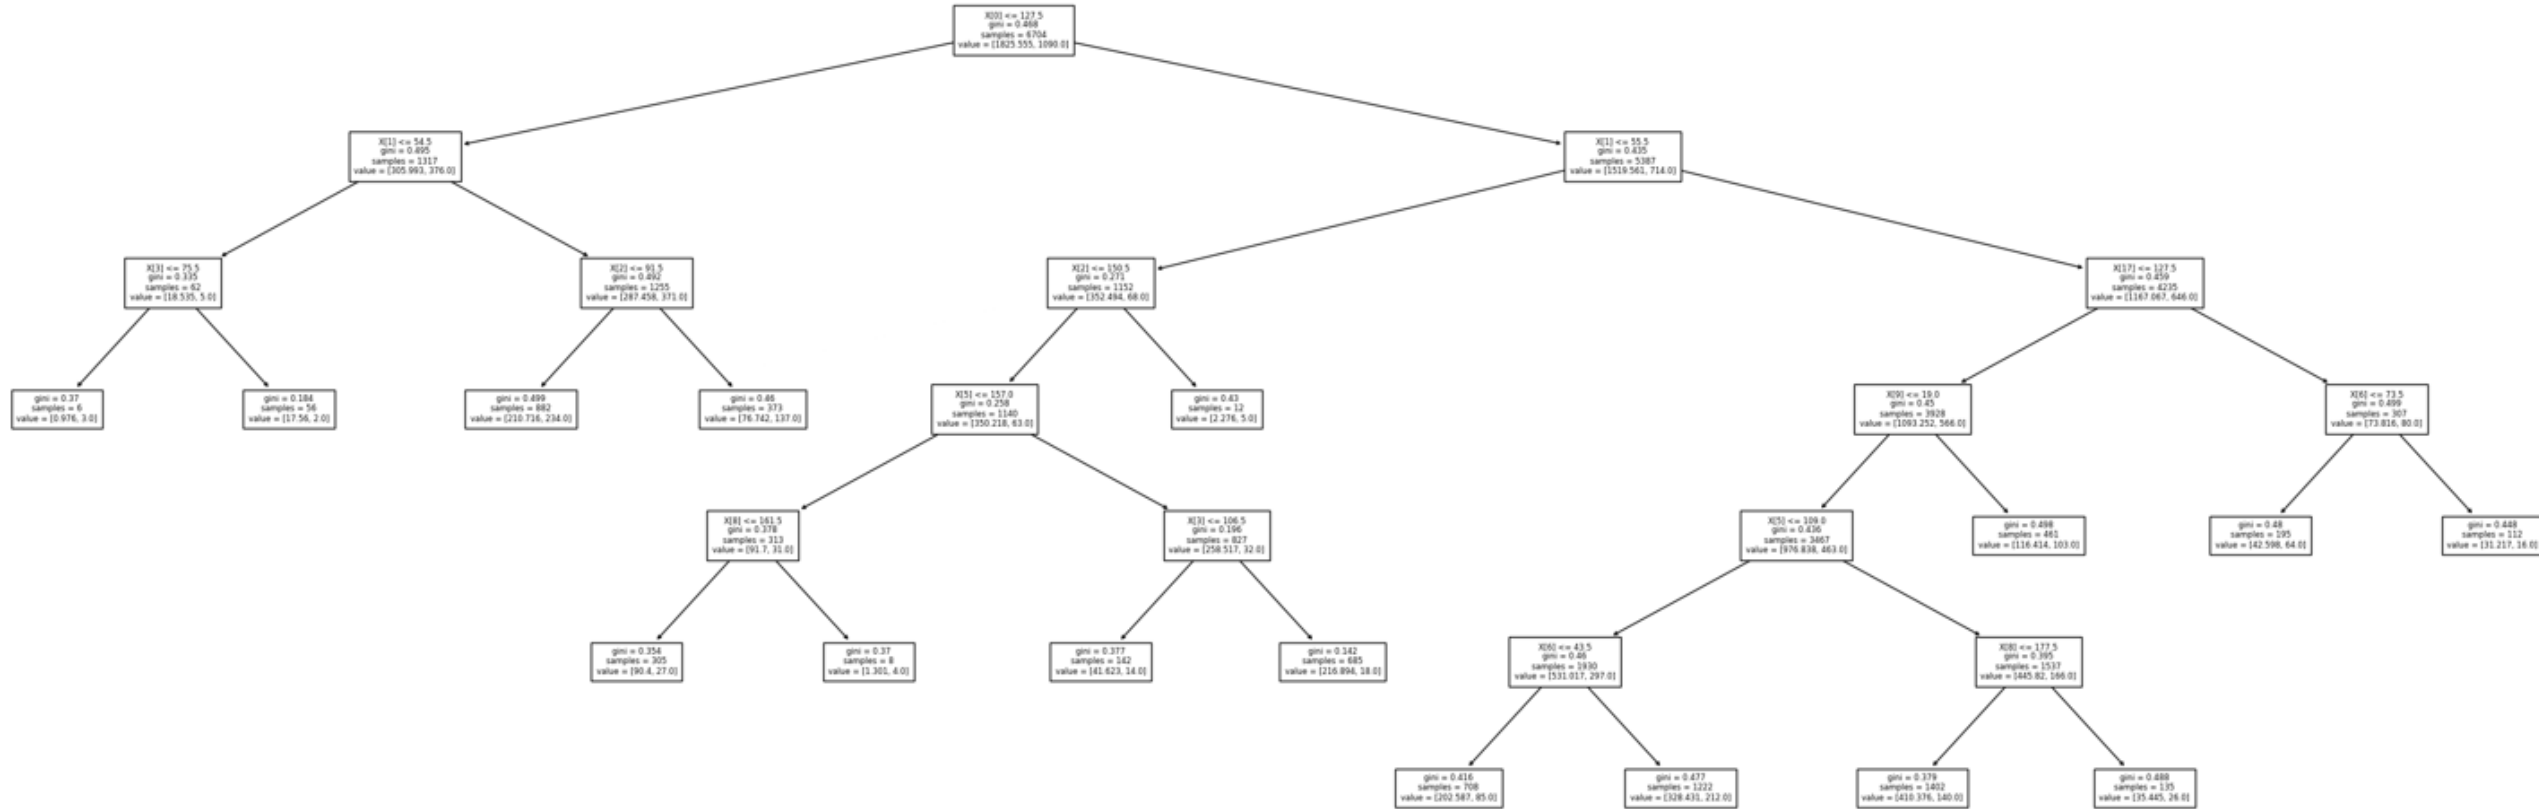
- Use the decision tree provided by scikit-learn
  - ***This work focused on the hardware implementation***
    - Fine-tuning and software analyses (exp: max-leaf-num, AUROC) are neglected
  - Confusion matrix on the test set (Before/ After converting to integers)

Ans\Pred	T	F
T	1970	461
F	284	159

Ans\Pred	T	F
T	1967	464
F	283	160

# Result Tree Structure

- Max Node Index: 30
- Max-depth: 6






# Outline

- Problem/ Dataset
- Software
- **Conversion**
- Hardware


# Observation

- Naïve thought
  - Core function is to “***move down 1 layer***”
- To move down 1 layer, we need 3 parameters:
  - The index of feature
  - The threshold
  - The index of the children



# Conversion Strategy

- Convert python code to Verilog hardware design
  - Parse all the parameters into 3 memories
  - Leaf nodes share the same slot



Node 0

Node 2  
(root  
node)

Feature Index	Threshold	Child
0	0	0
0	0	0
1	5	0
3	225	0
0	0	0
0	0	3
0	0	0
0	0	1

# Feature/Node Pruning

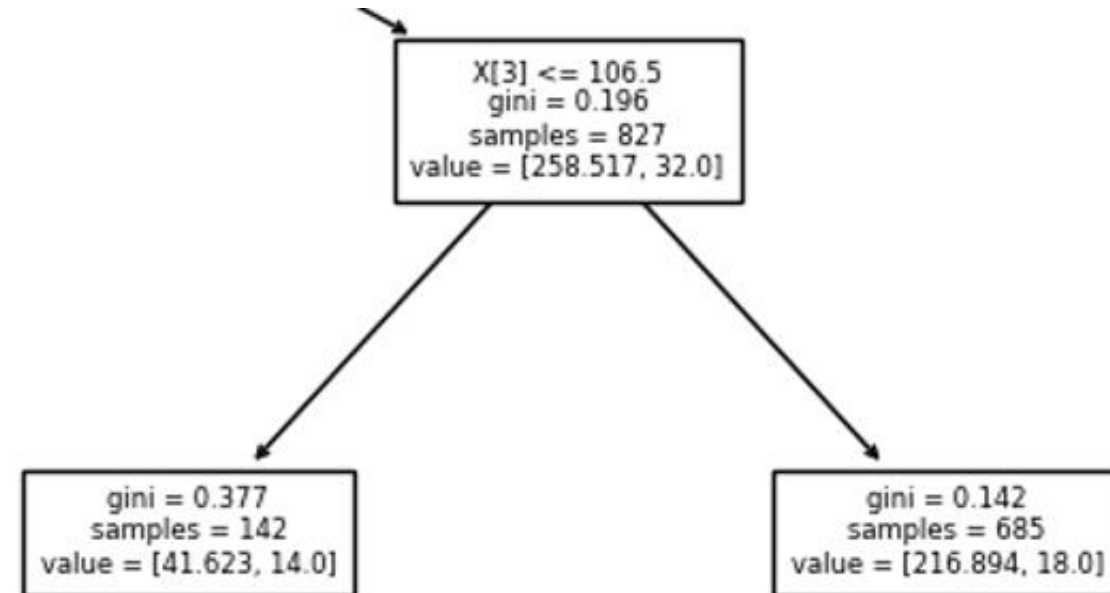
- The decision tree generated from scikit-learn can be further optimized
- Not all 18 features are used
  - Prune the unused ones. Re-index the features left.
- Not all 31 node indexes are occupied
  - Leaf nodes do not need a unique index!
  - Prune the unused ones. Re-index the nodes.

# Redundant Splits Pruning

- Reason: Decision tree training target
  - Minimizing the weighted entropy sum after split
    - *Does not guarantee the split produces different predictions!*
  - Redundant splits are sometimes generated
  - Recursively prune the redundant splits

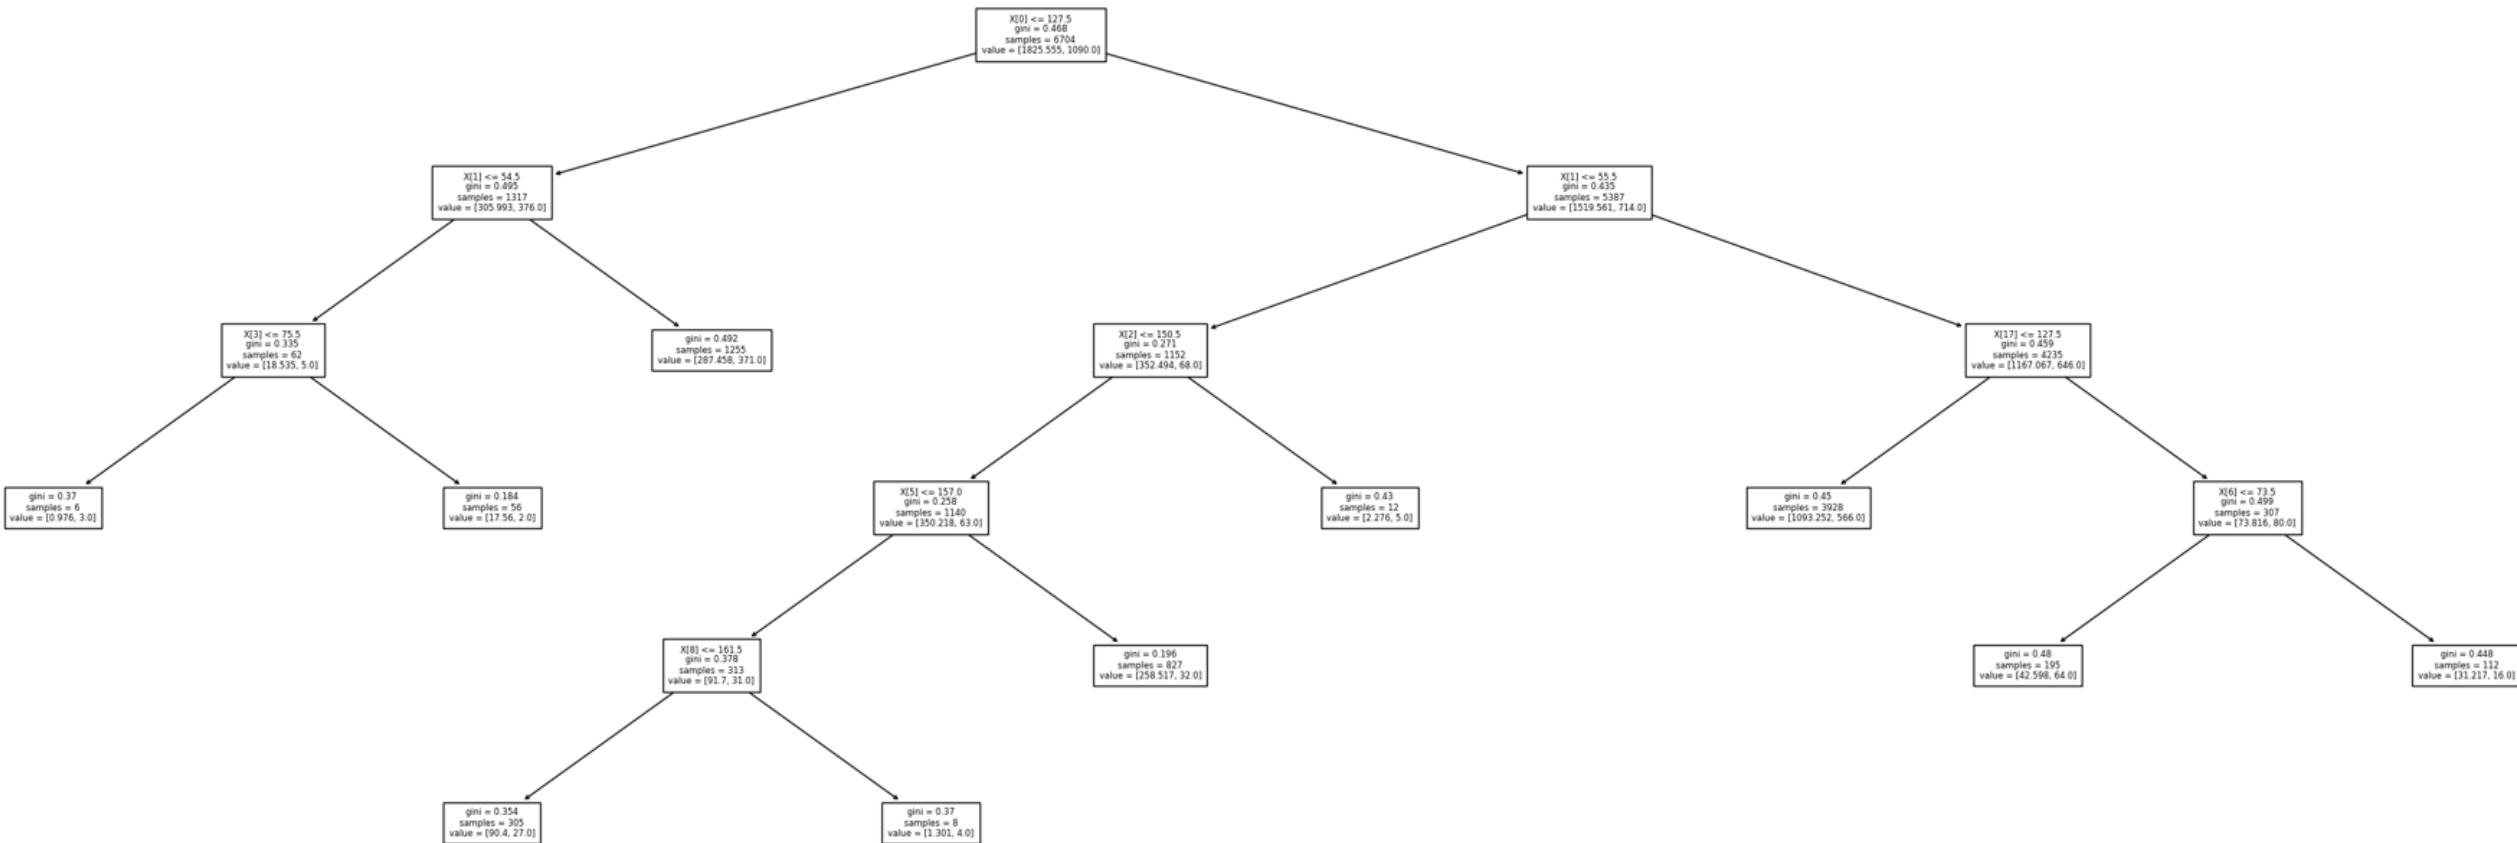
$$0.377 * 142 + 0.142 * 685 < 0.196 * 827$$

But both leaf nodes predict class 0!



# Final Tree Structure

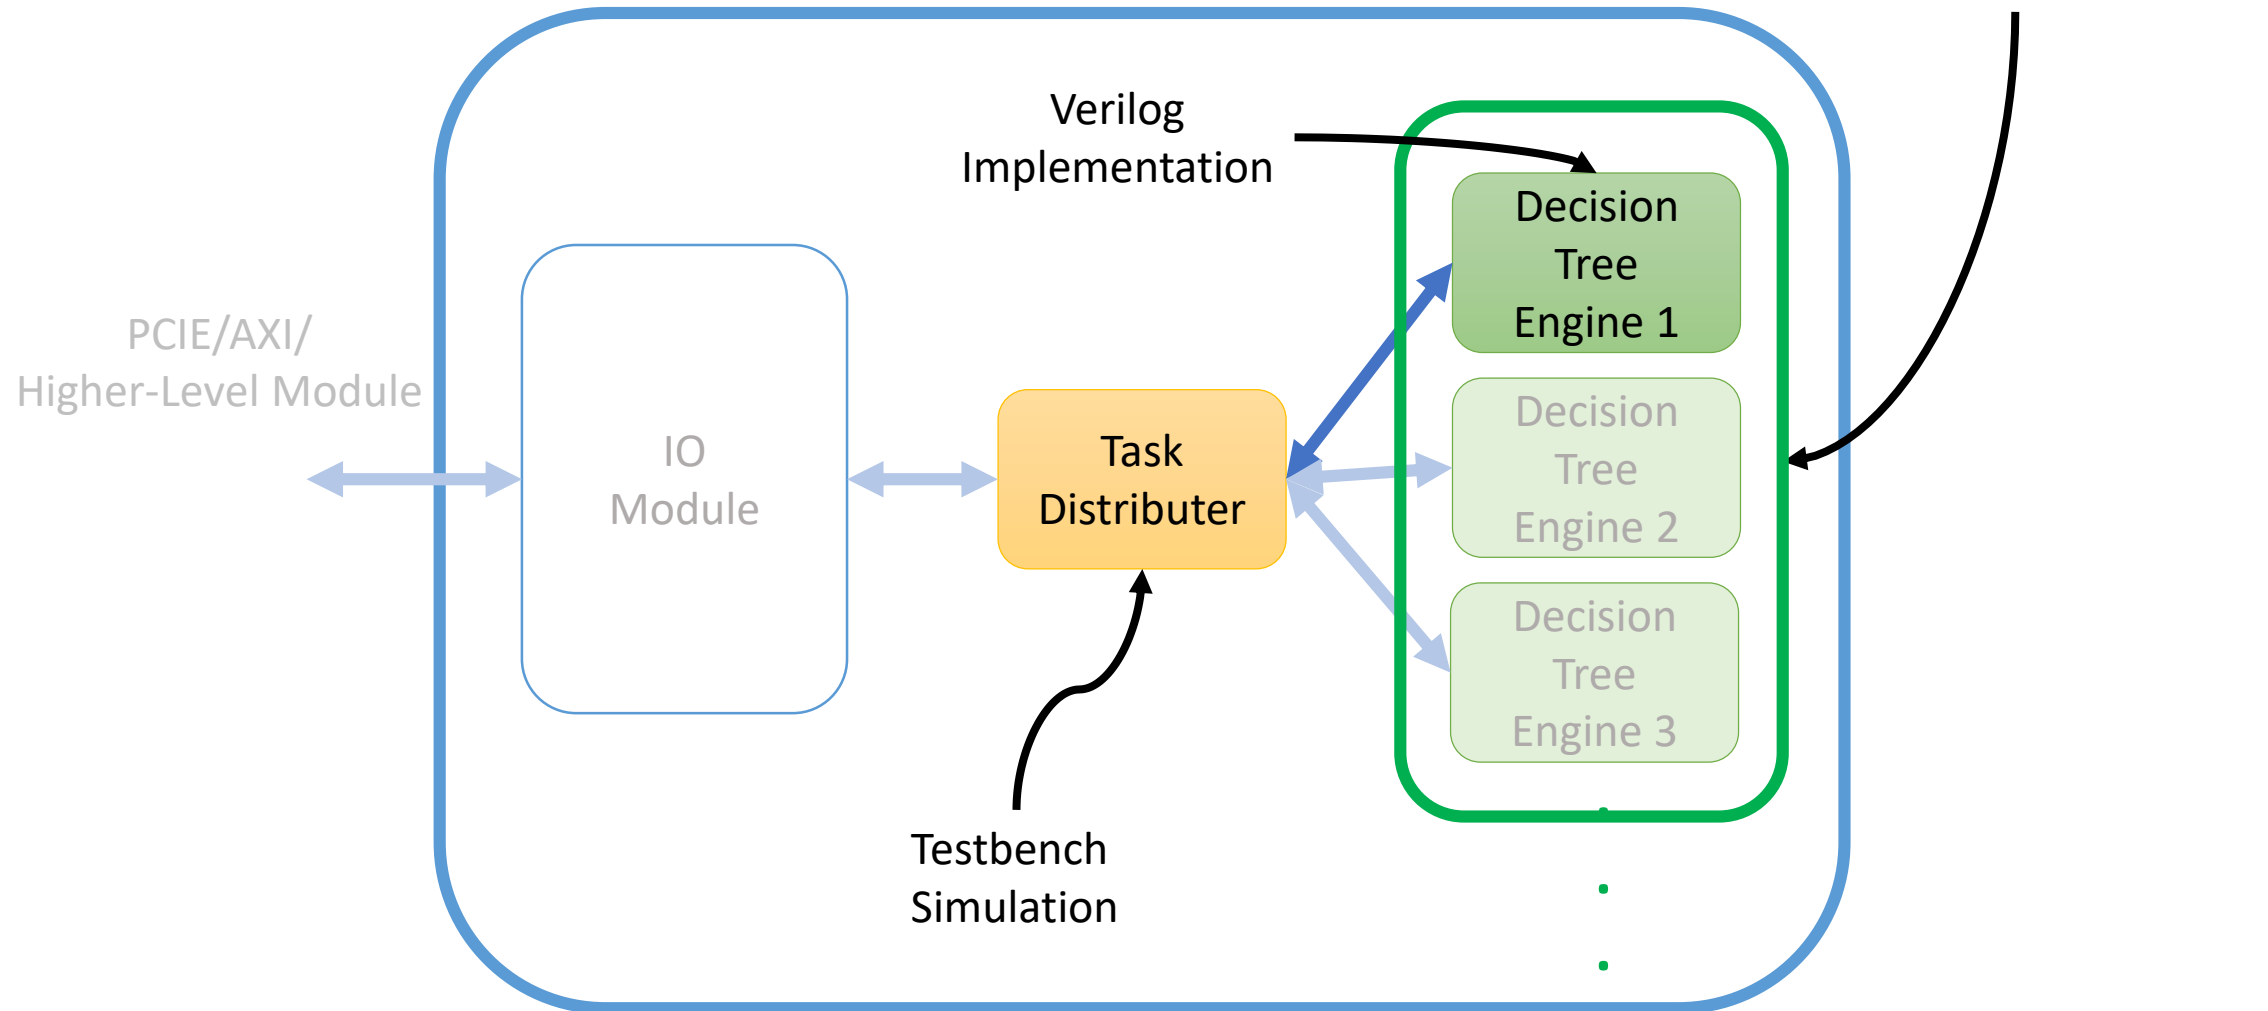
- Max node index: 10
- Max-depth: 5 layers



# Outline

- Problem/ Dataset
- Software
- Conversion
- Hardware

# Top-Level View / Project Scope



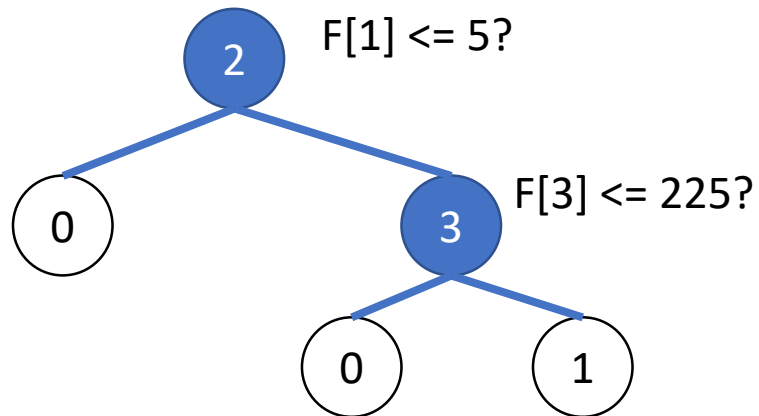


# Architecture – Core Function

- 3-stage pipeline to move 1 layer down in the tree
  - 1<sup>st</sup>:
    - Retrieve the feature index by the node index
  - 2<sup>nd</sup>:
    - Select the feature using feature index
    - Retrieve the threshold used
  - 3<sup>rd</sup>:
    - Comparing the selected feature with the threshold
    - Retrieve index of the next node

# Architecture – Core Function

- Parameter storage option
  - SRAM: Better generalizability for larger trees!
  - Flip-Flop array



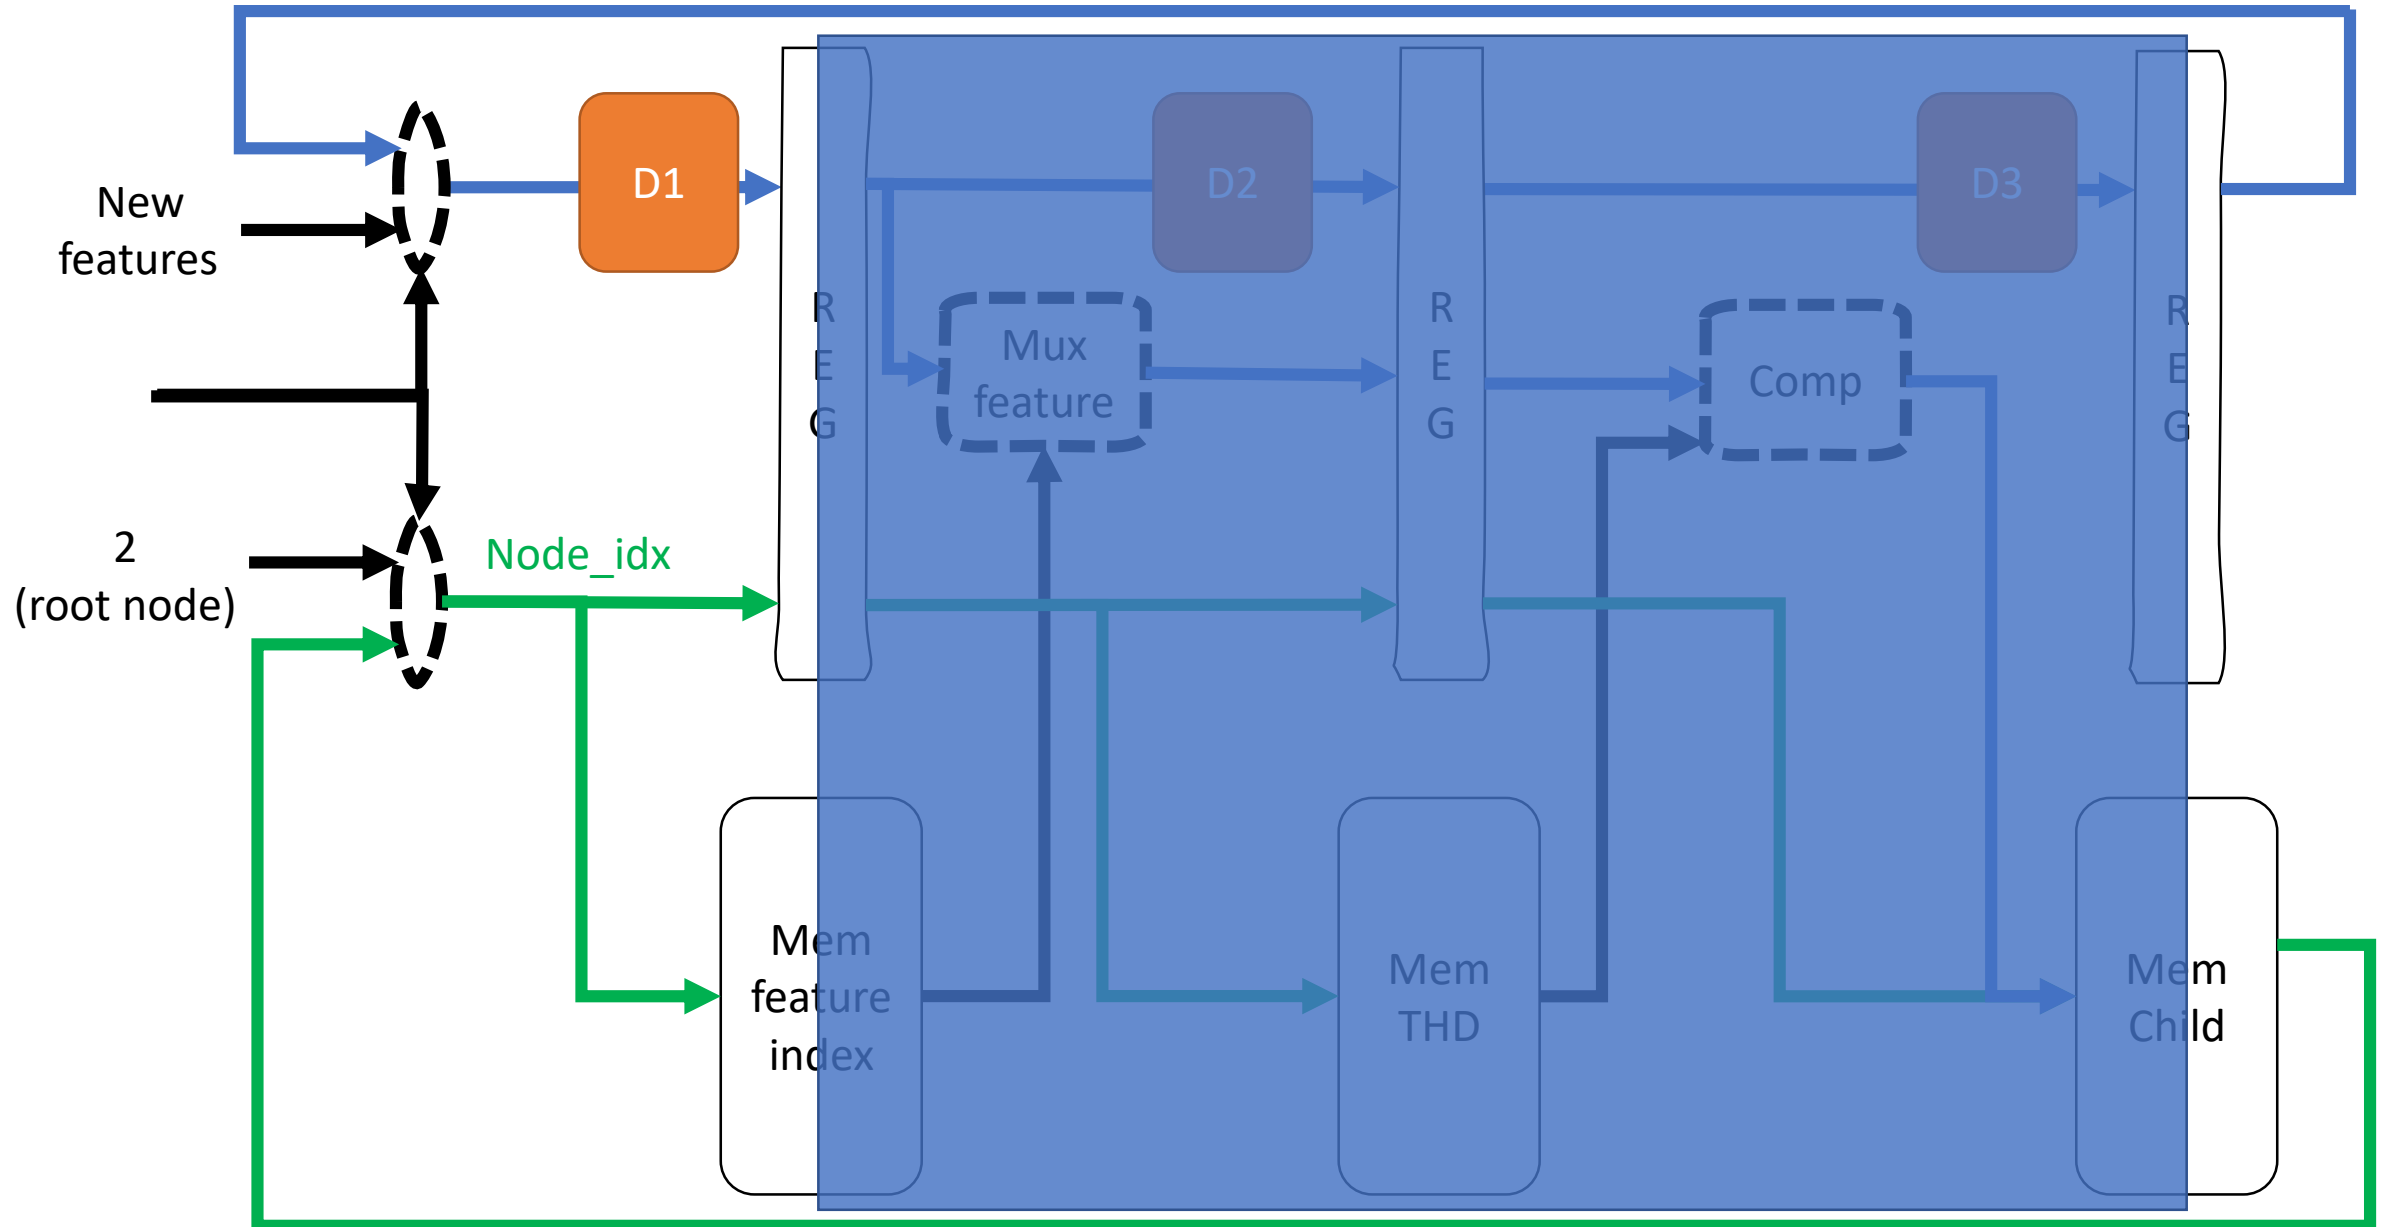
Node 0

Node 2  
(root  
node)

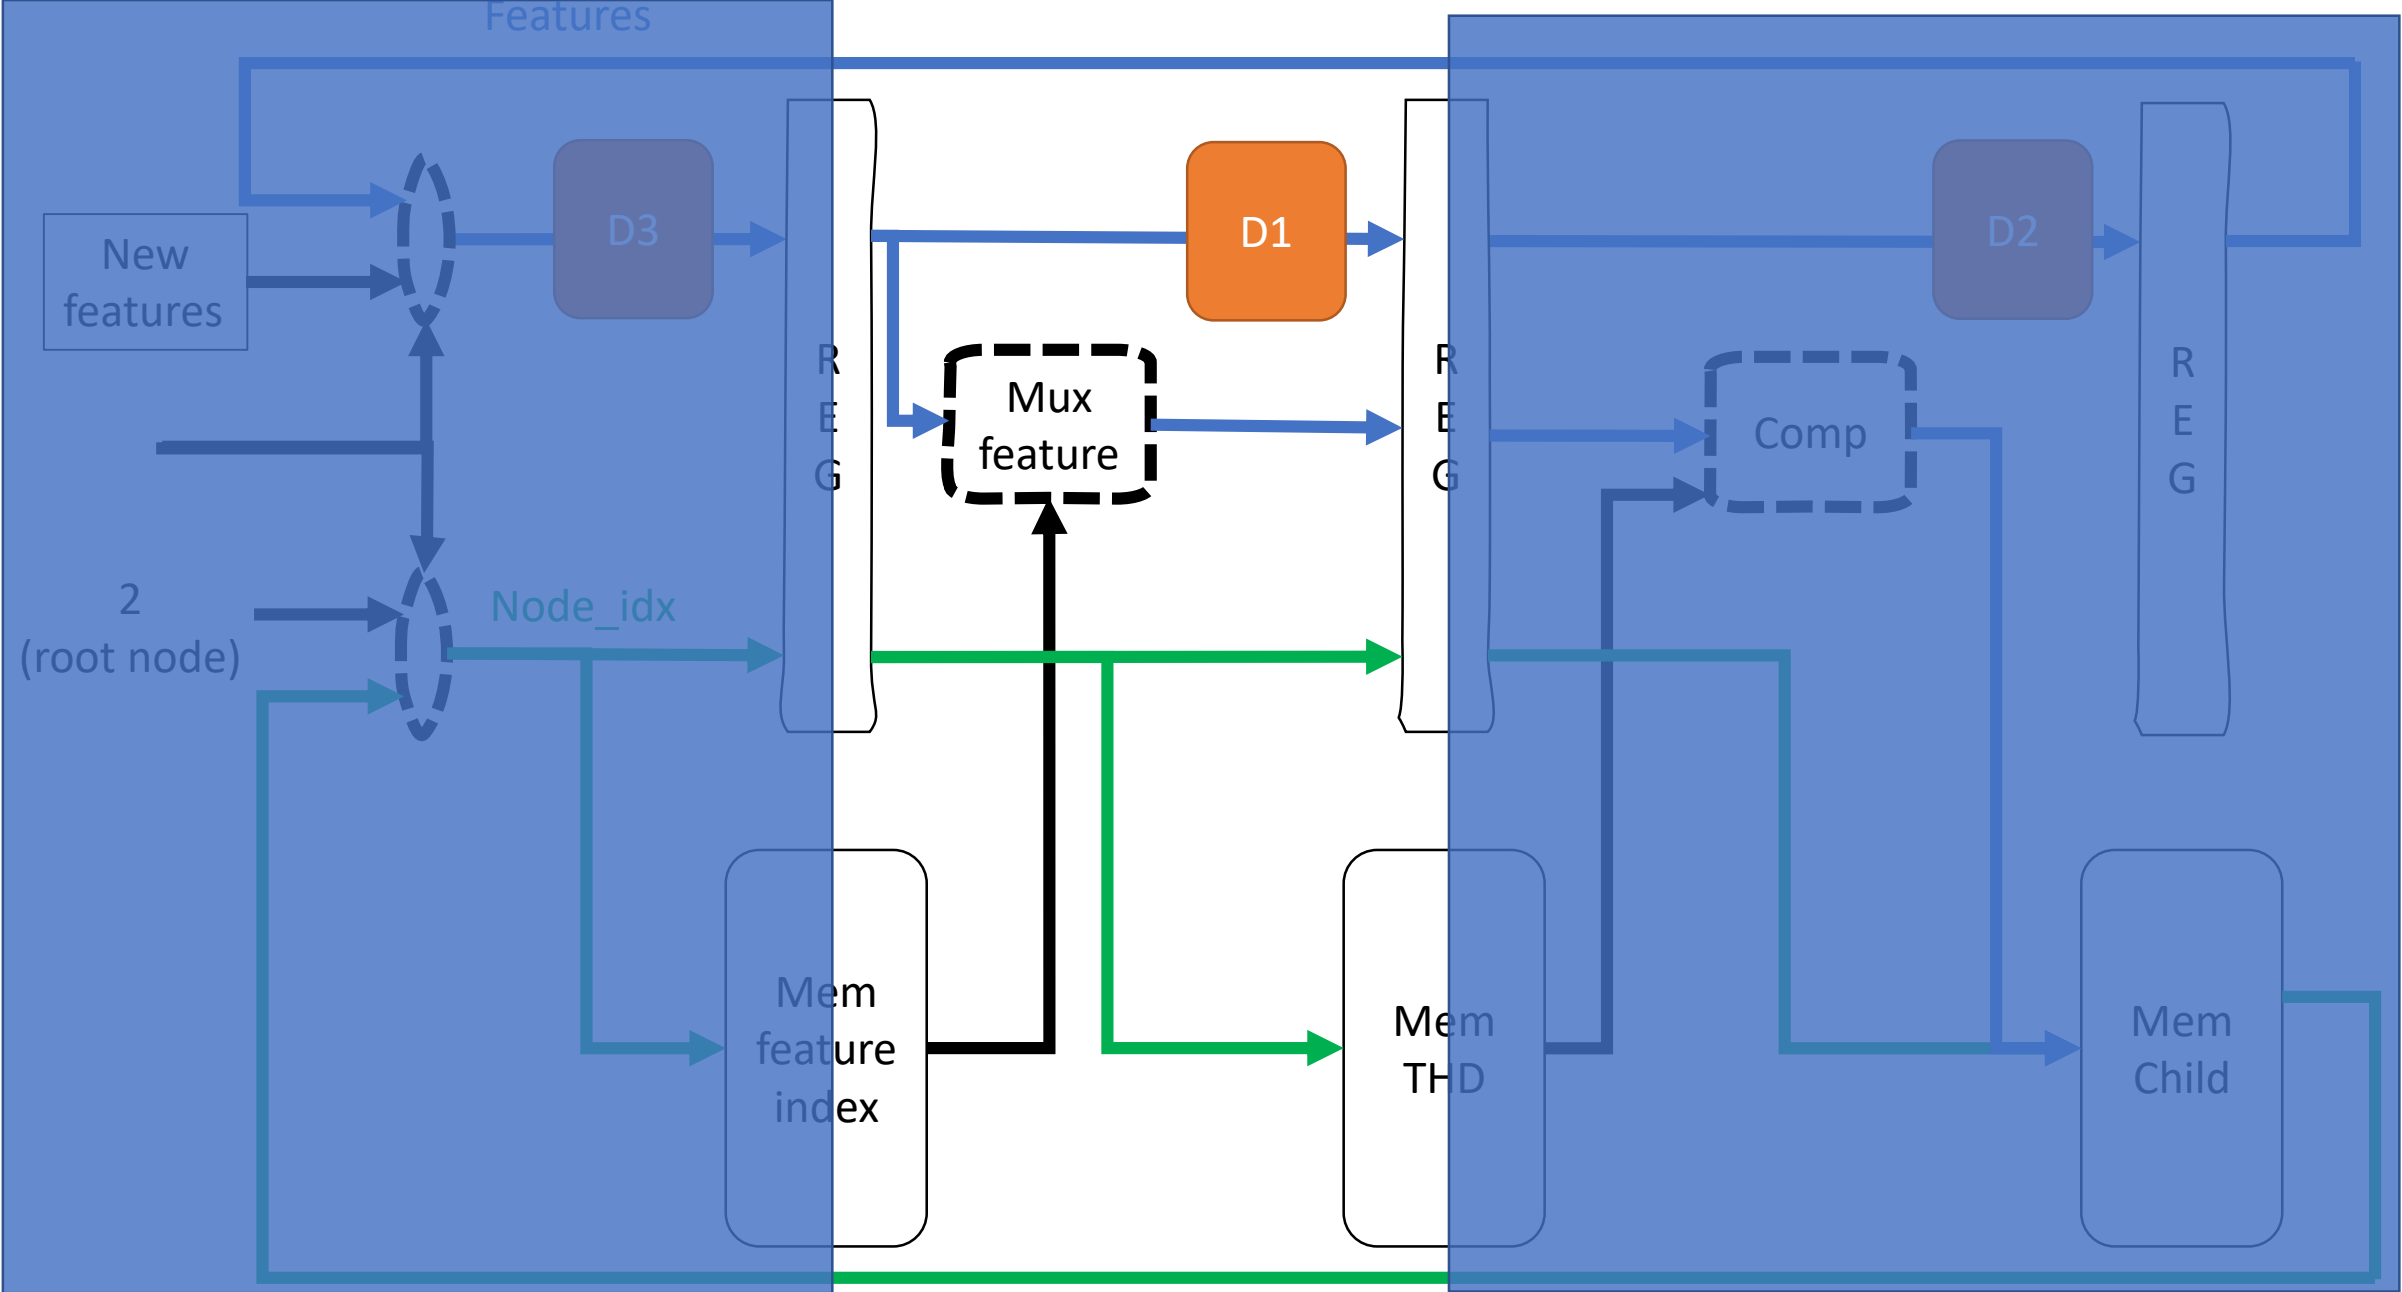
Feature Index	Threshold	Child
0	0	0
0	0	0
1	5	0
3	225	0
0	0	0
0	0	3
0	0	0
0	0	1

# Stage 1

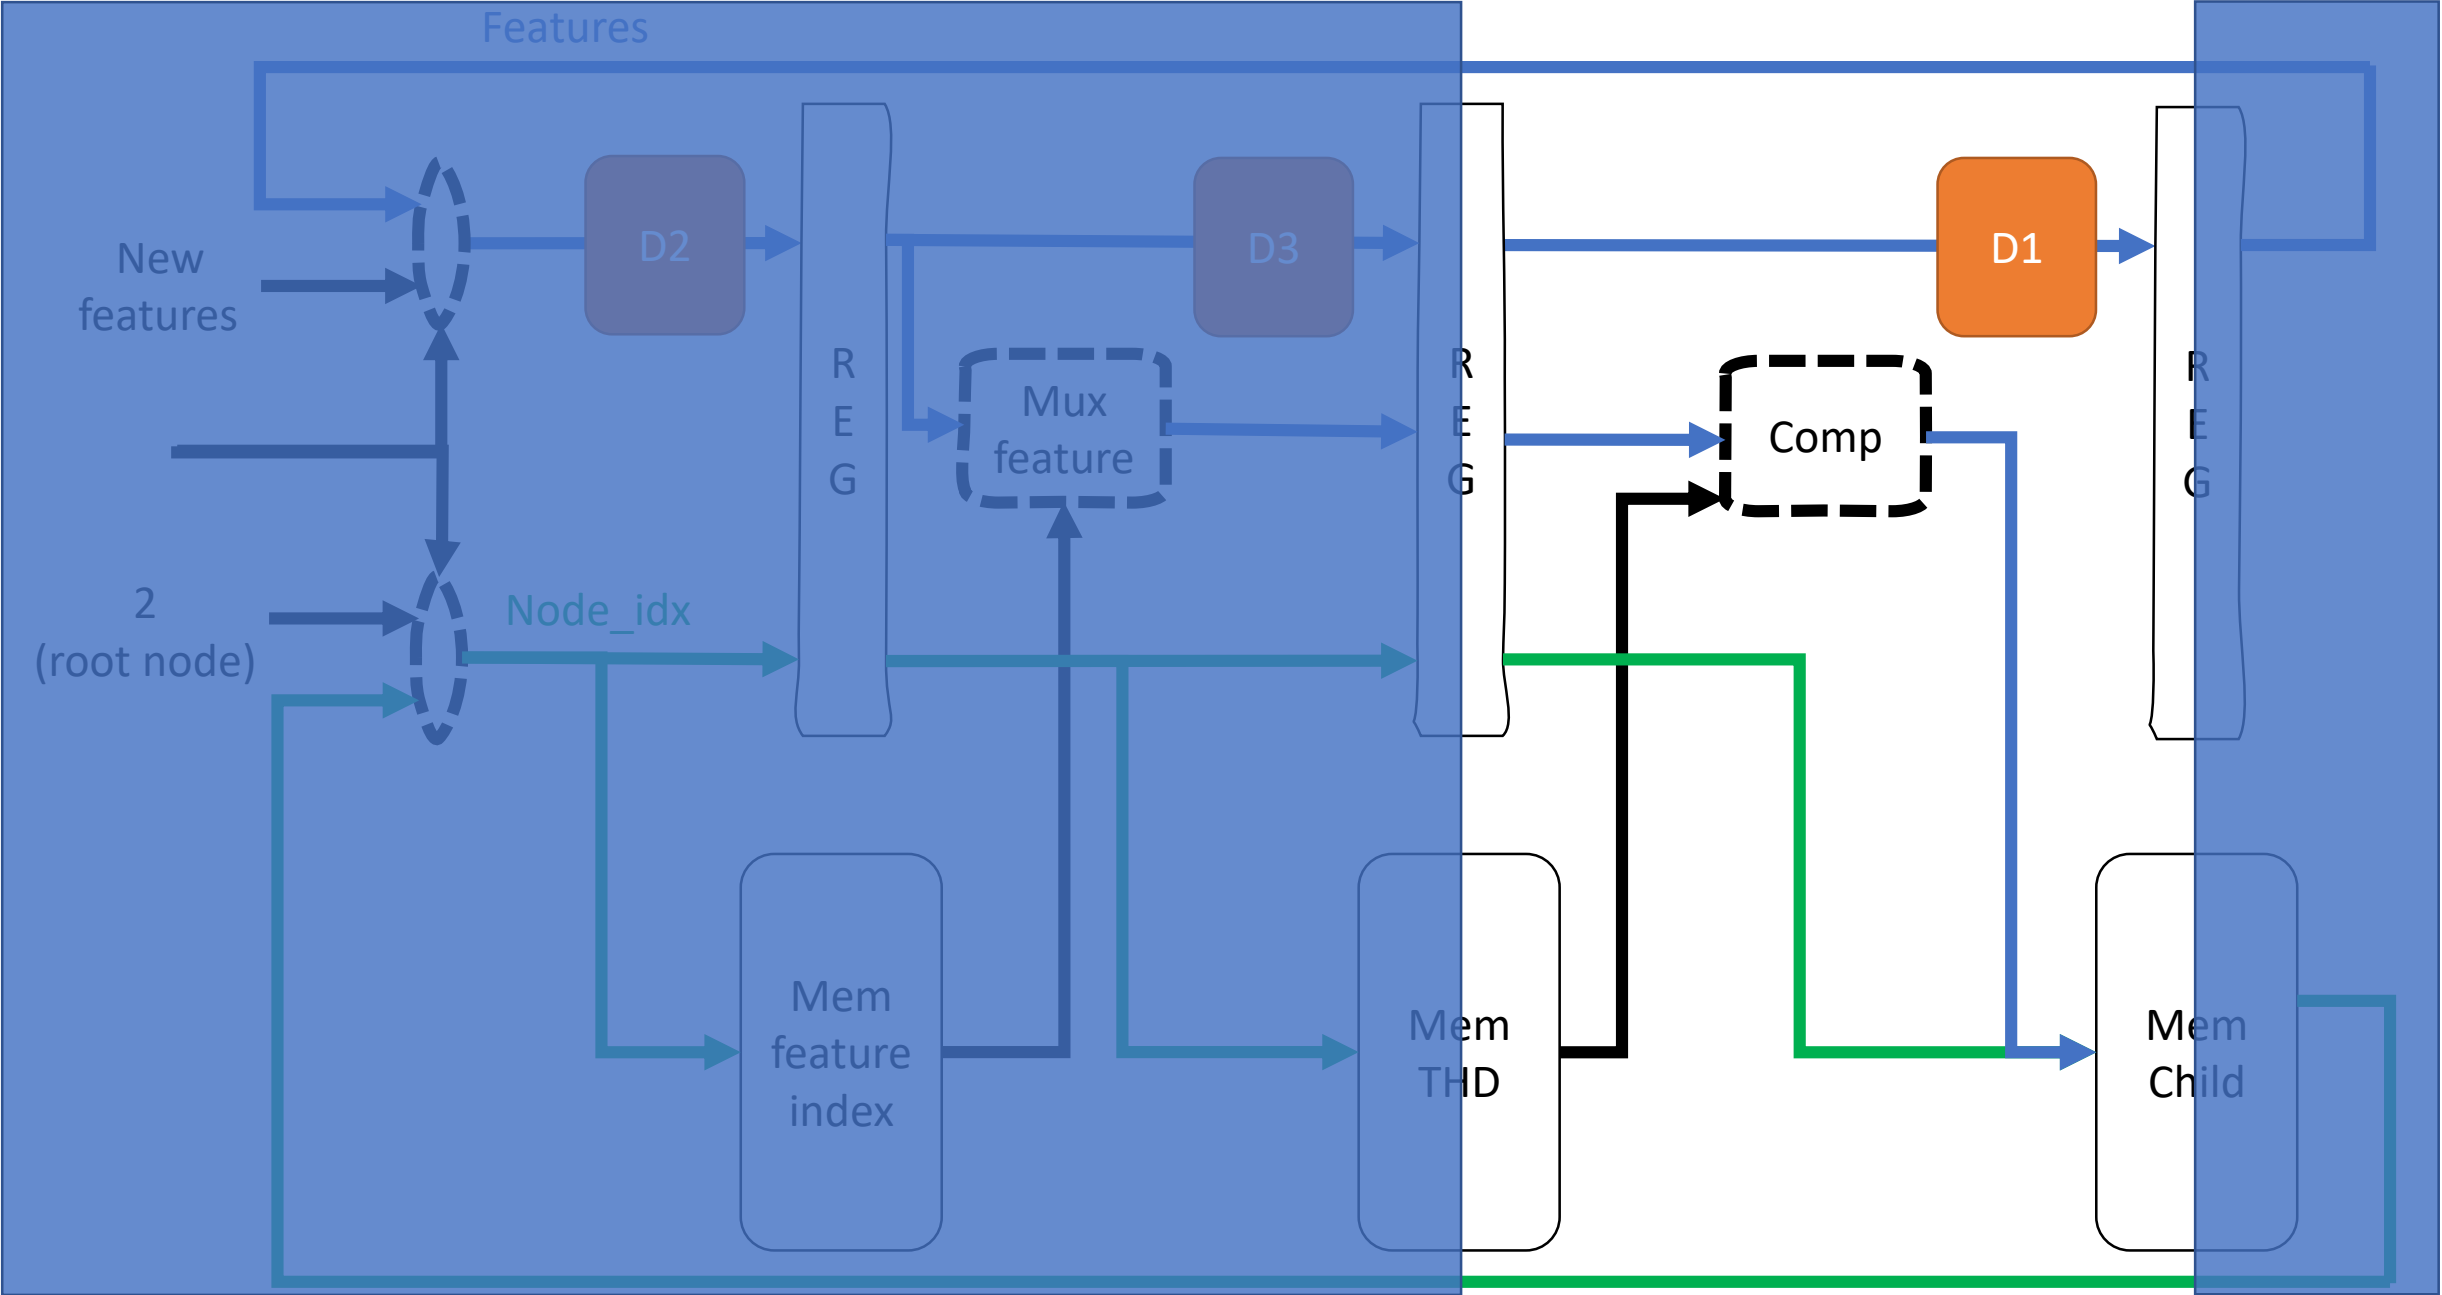
Features



**Stage 2**

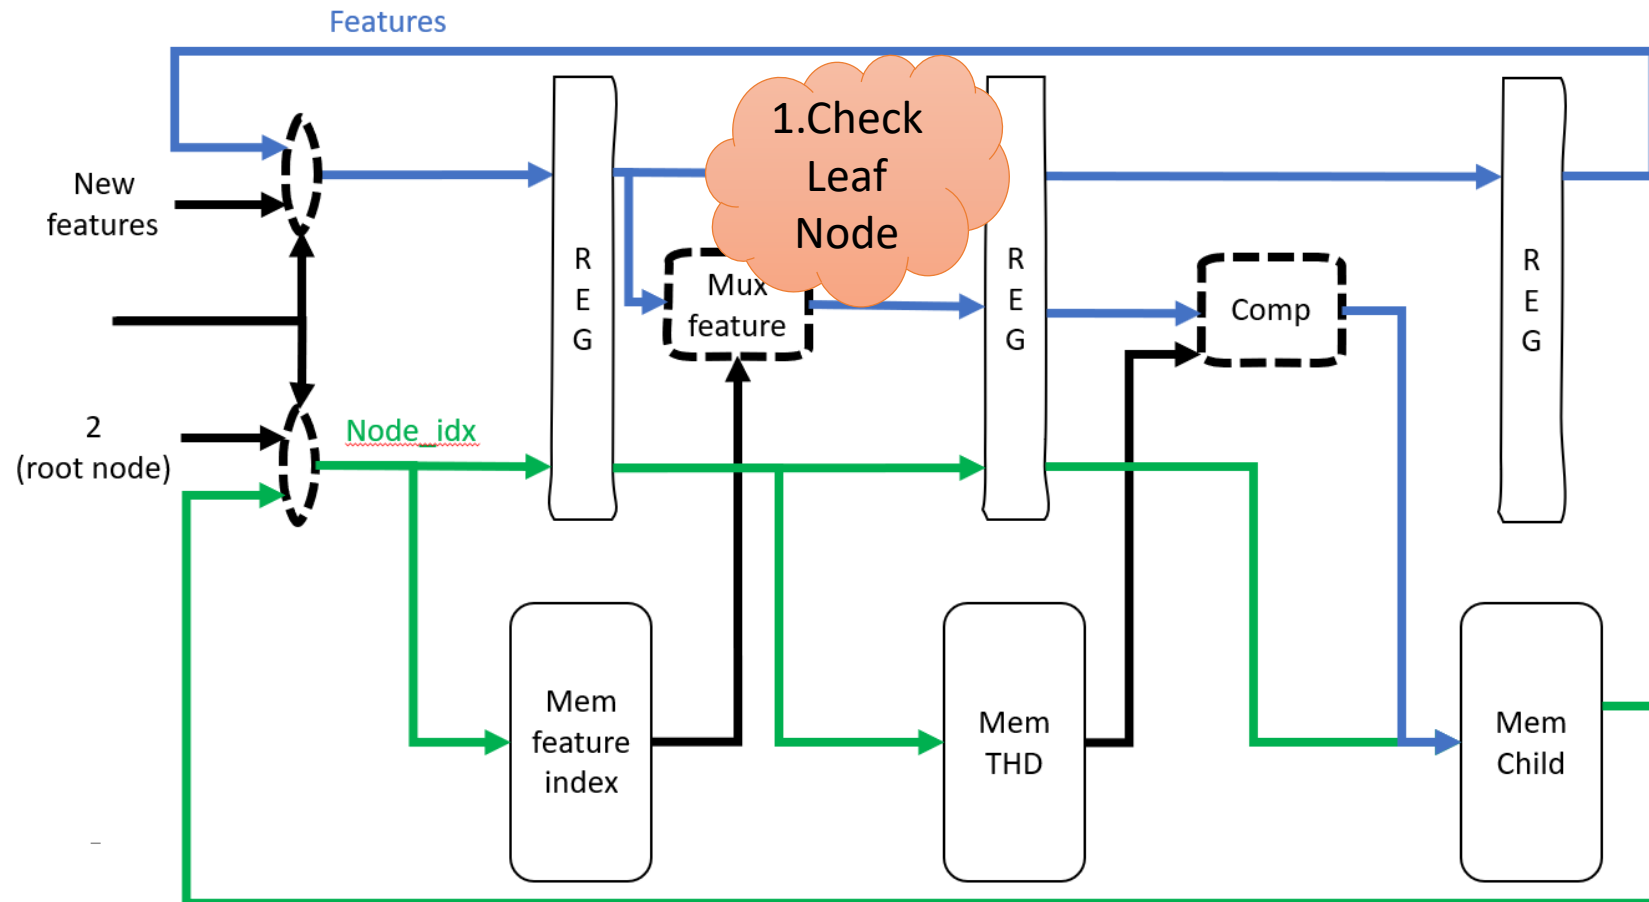


**Stage 3**



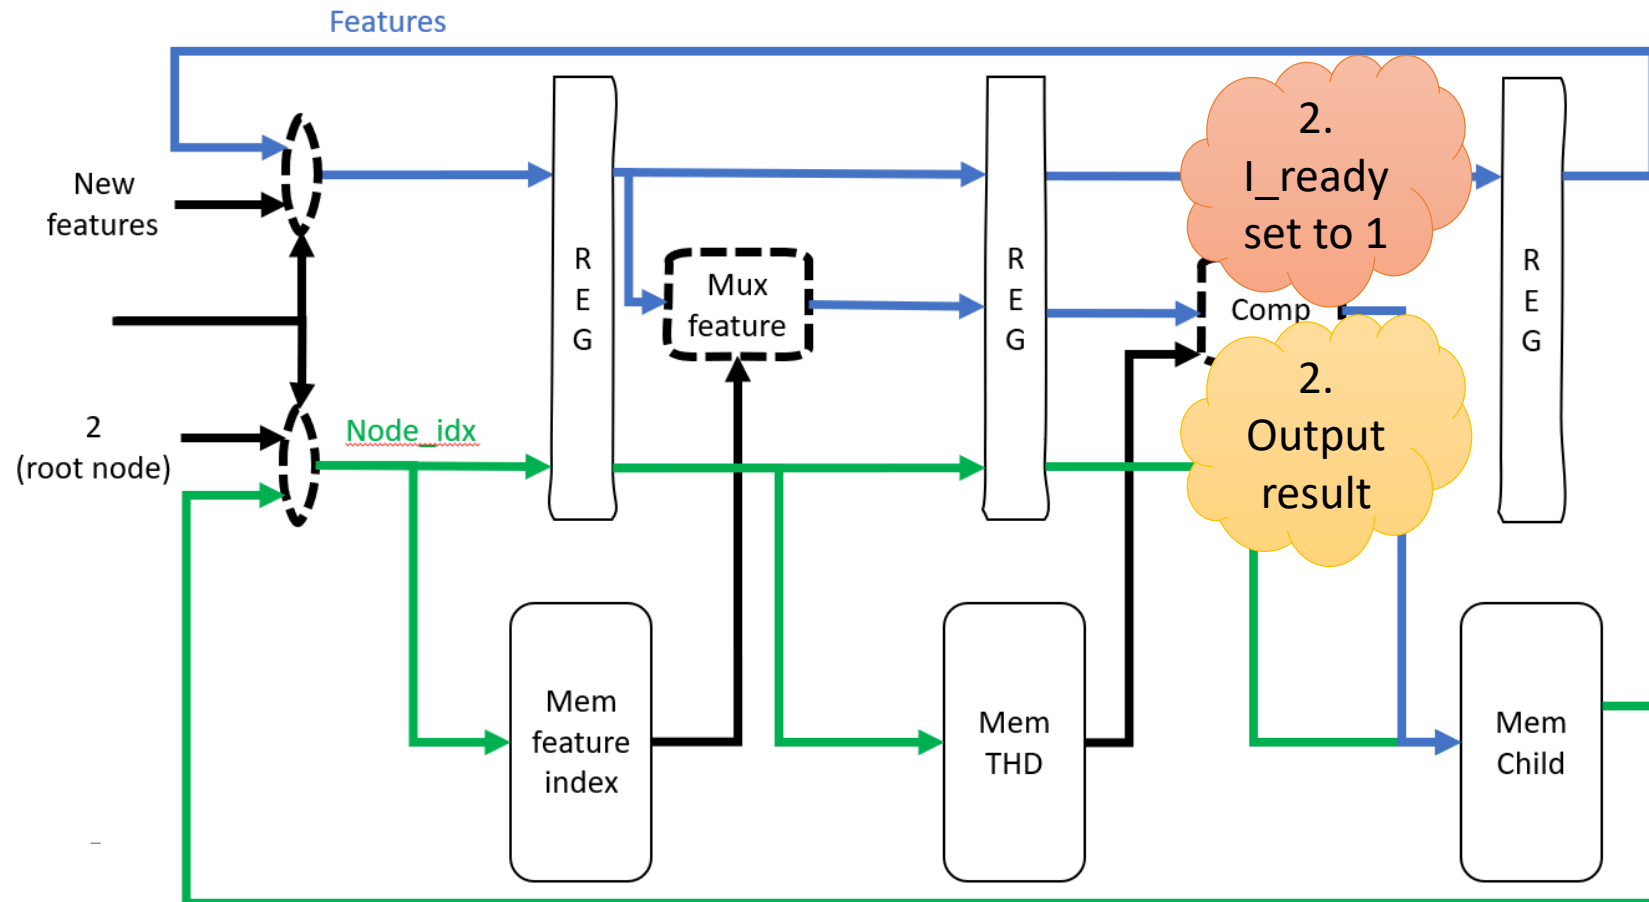
# Architecture – A deeper look

- Input / Output
  - Assume the task distributor takes 1 cycle to respond



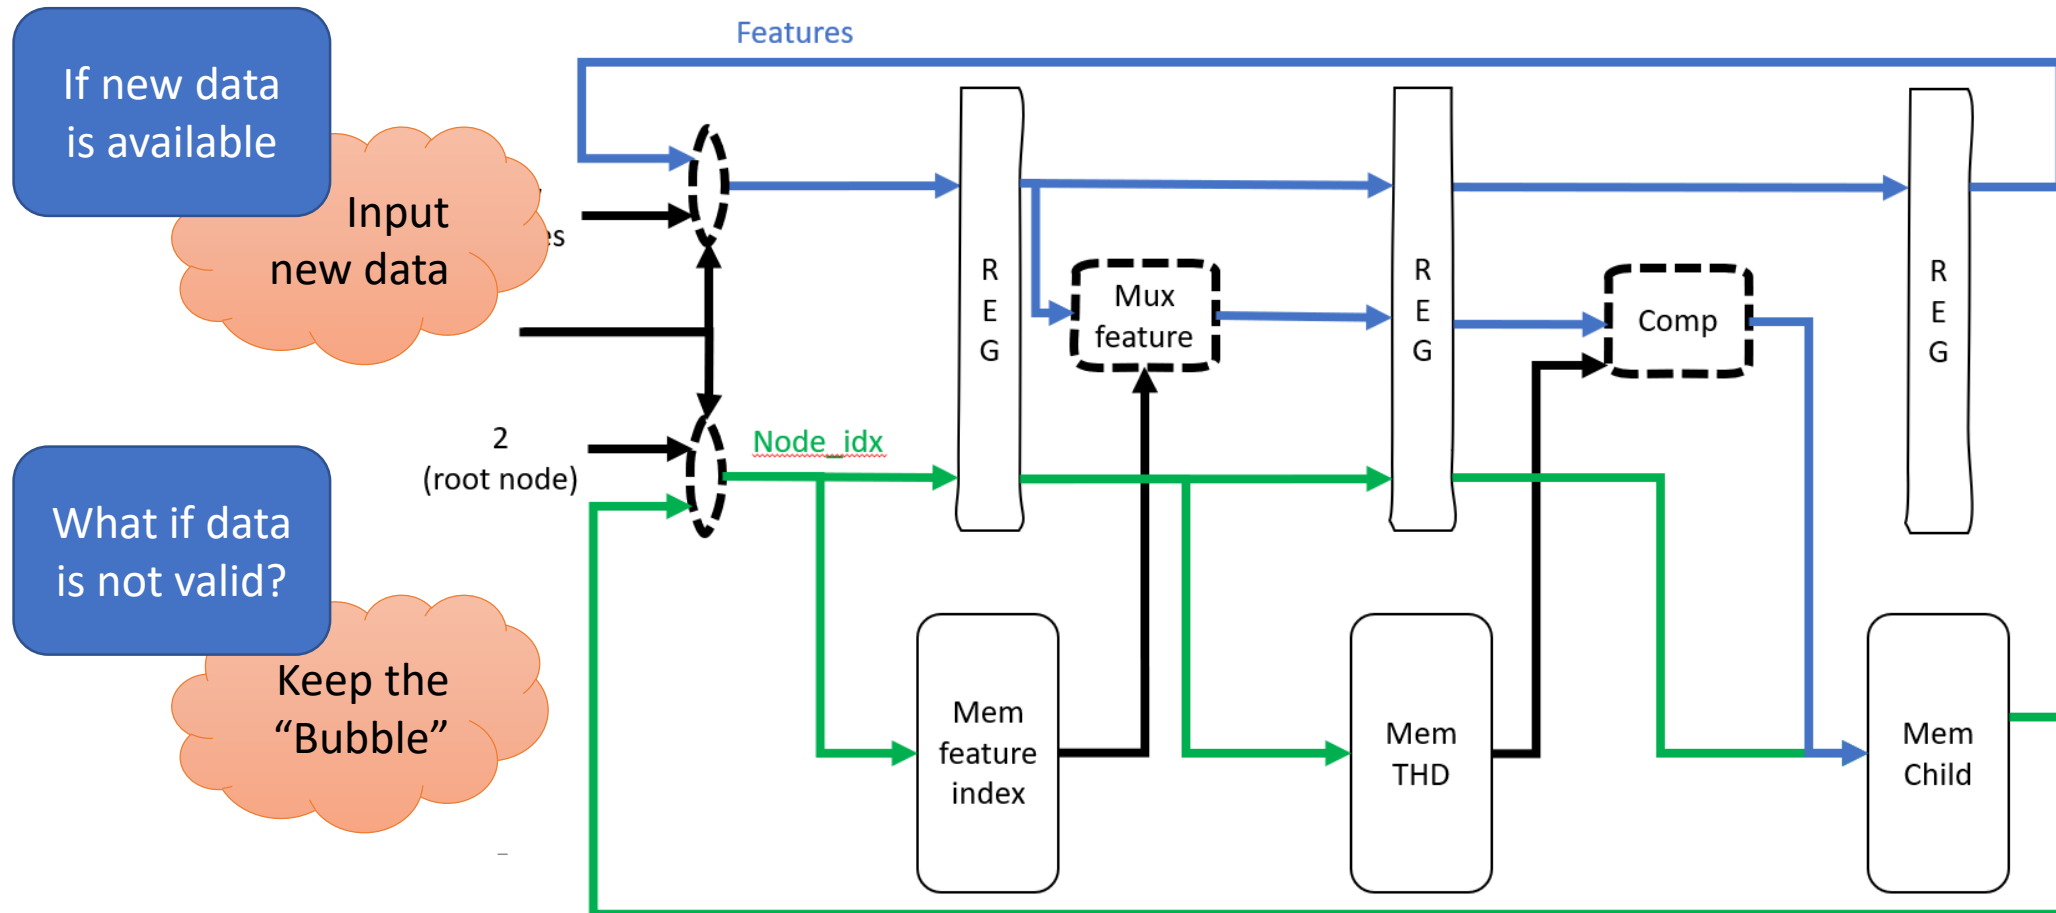
# Architecture – A deeper look

- Input / Output
  - Assume the task distributor takes 1 cycle to respond

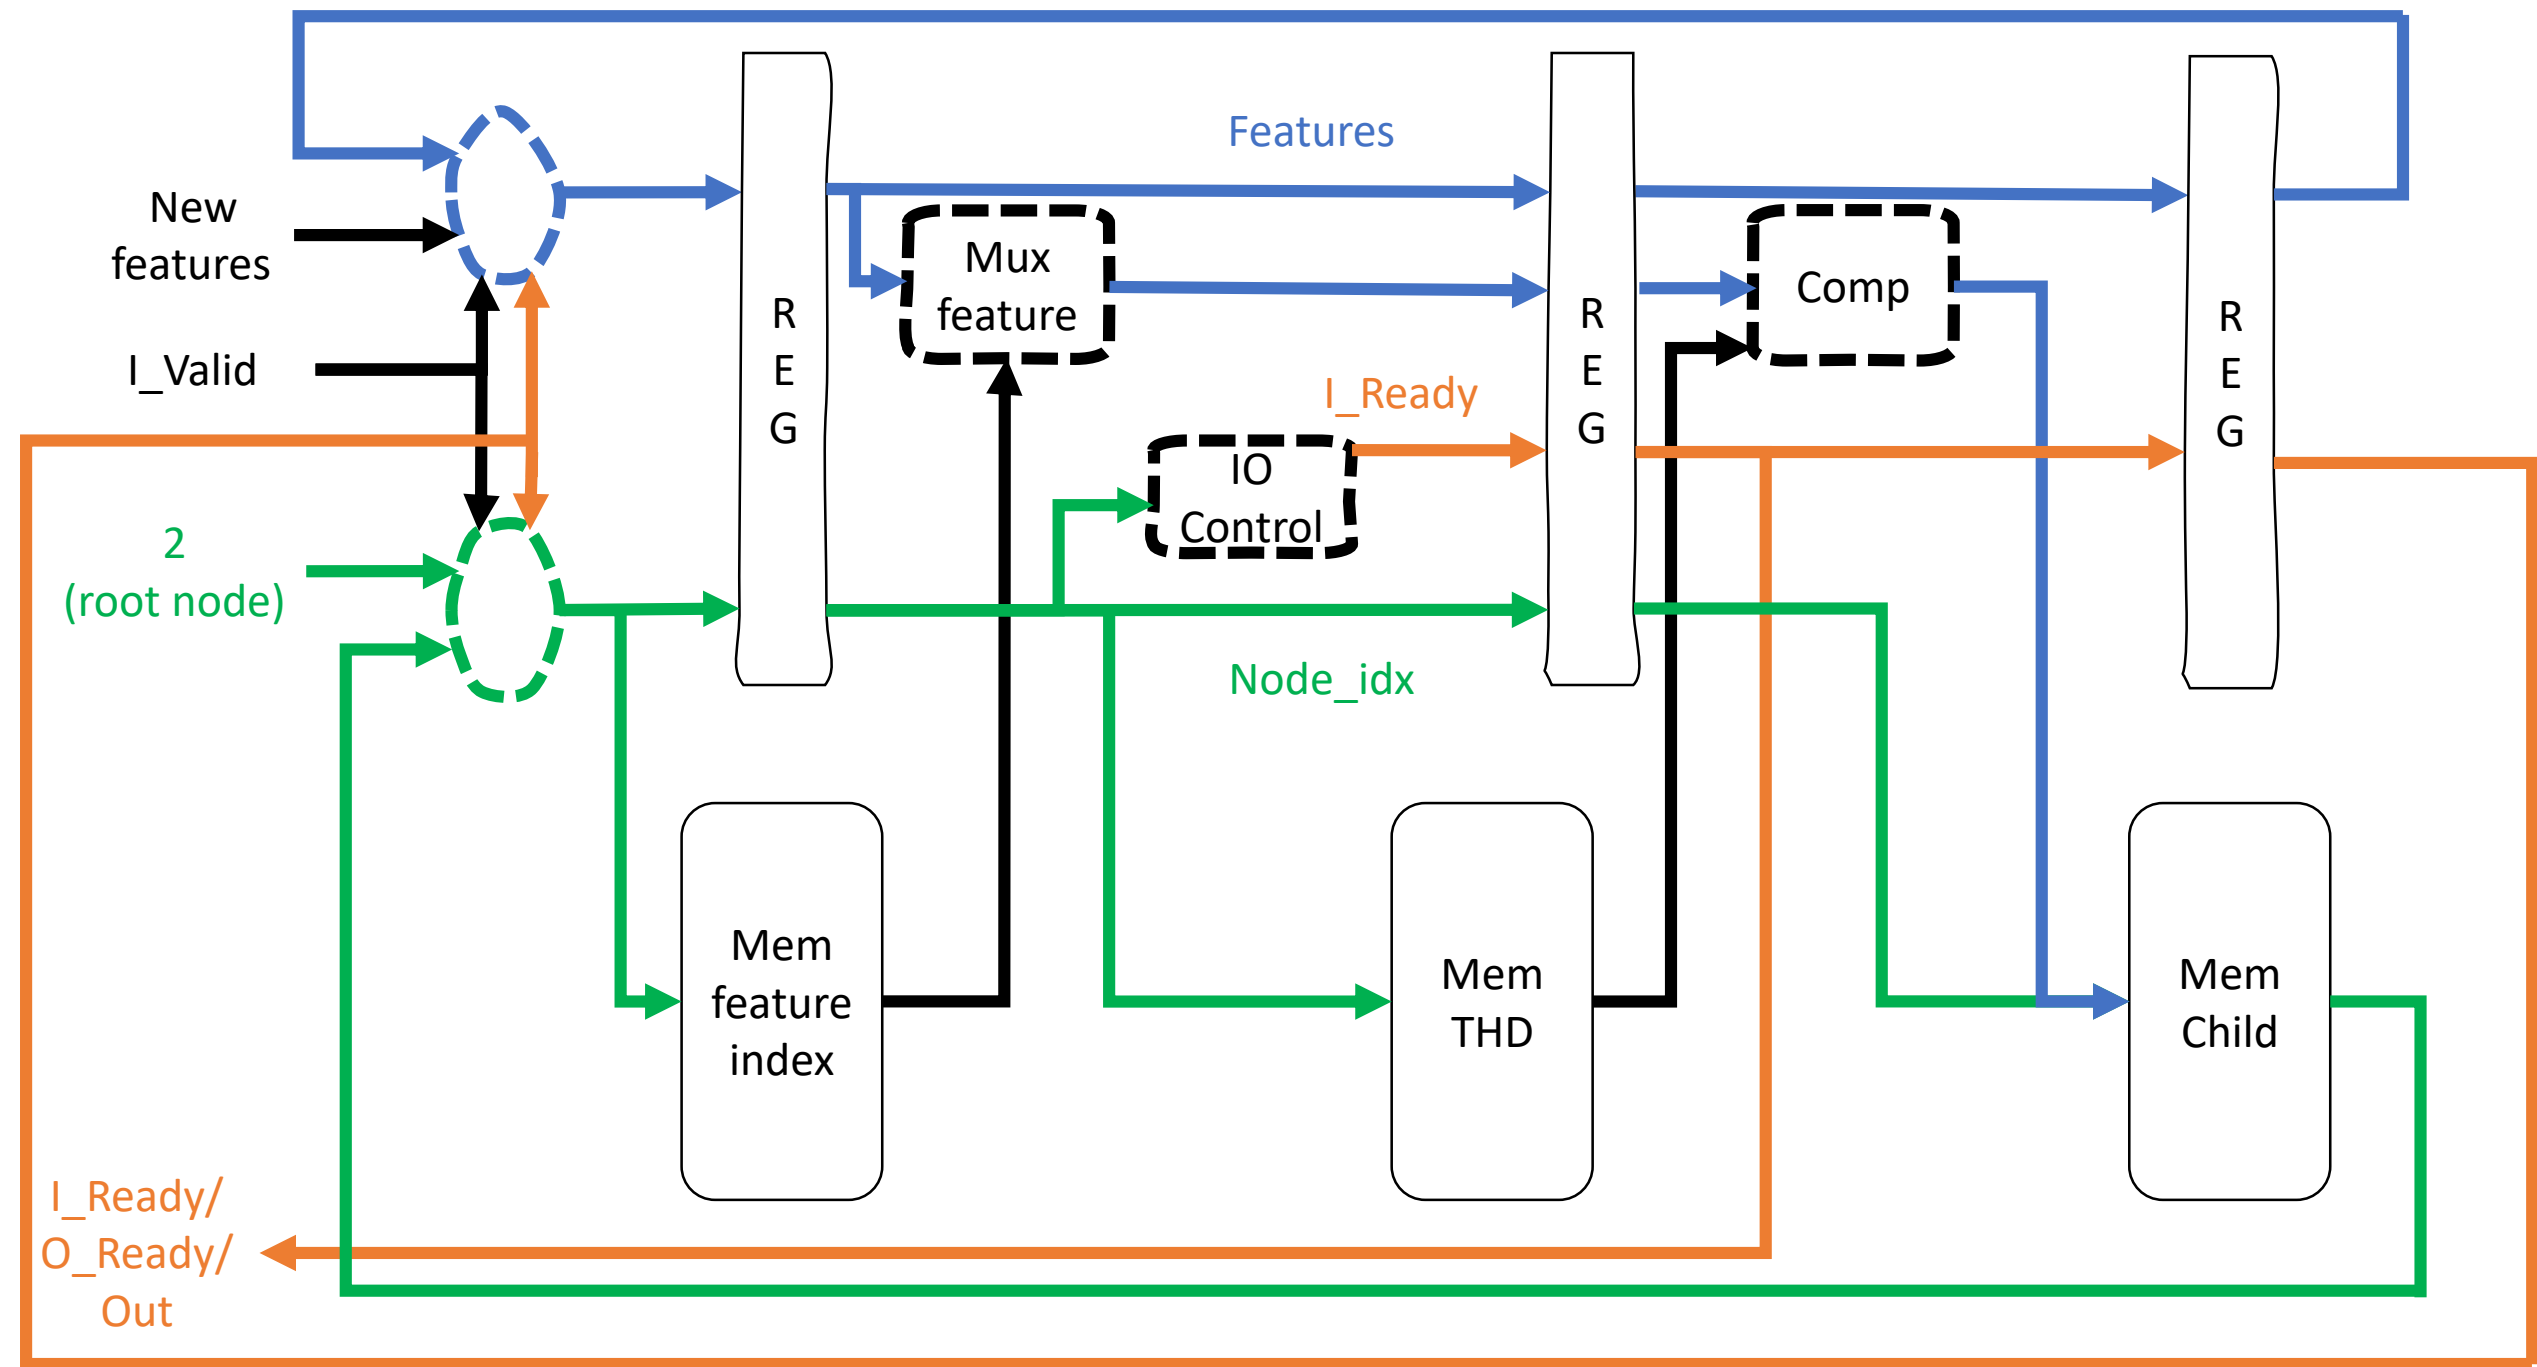


# Architecture – A deeper look

- Input / Output
  - Assume the task distributor takes 1 cycle to respond

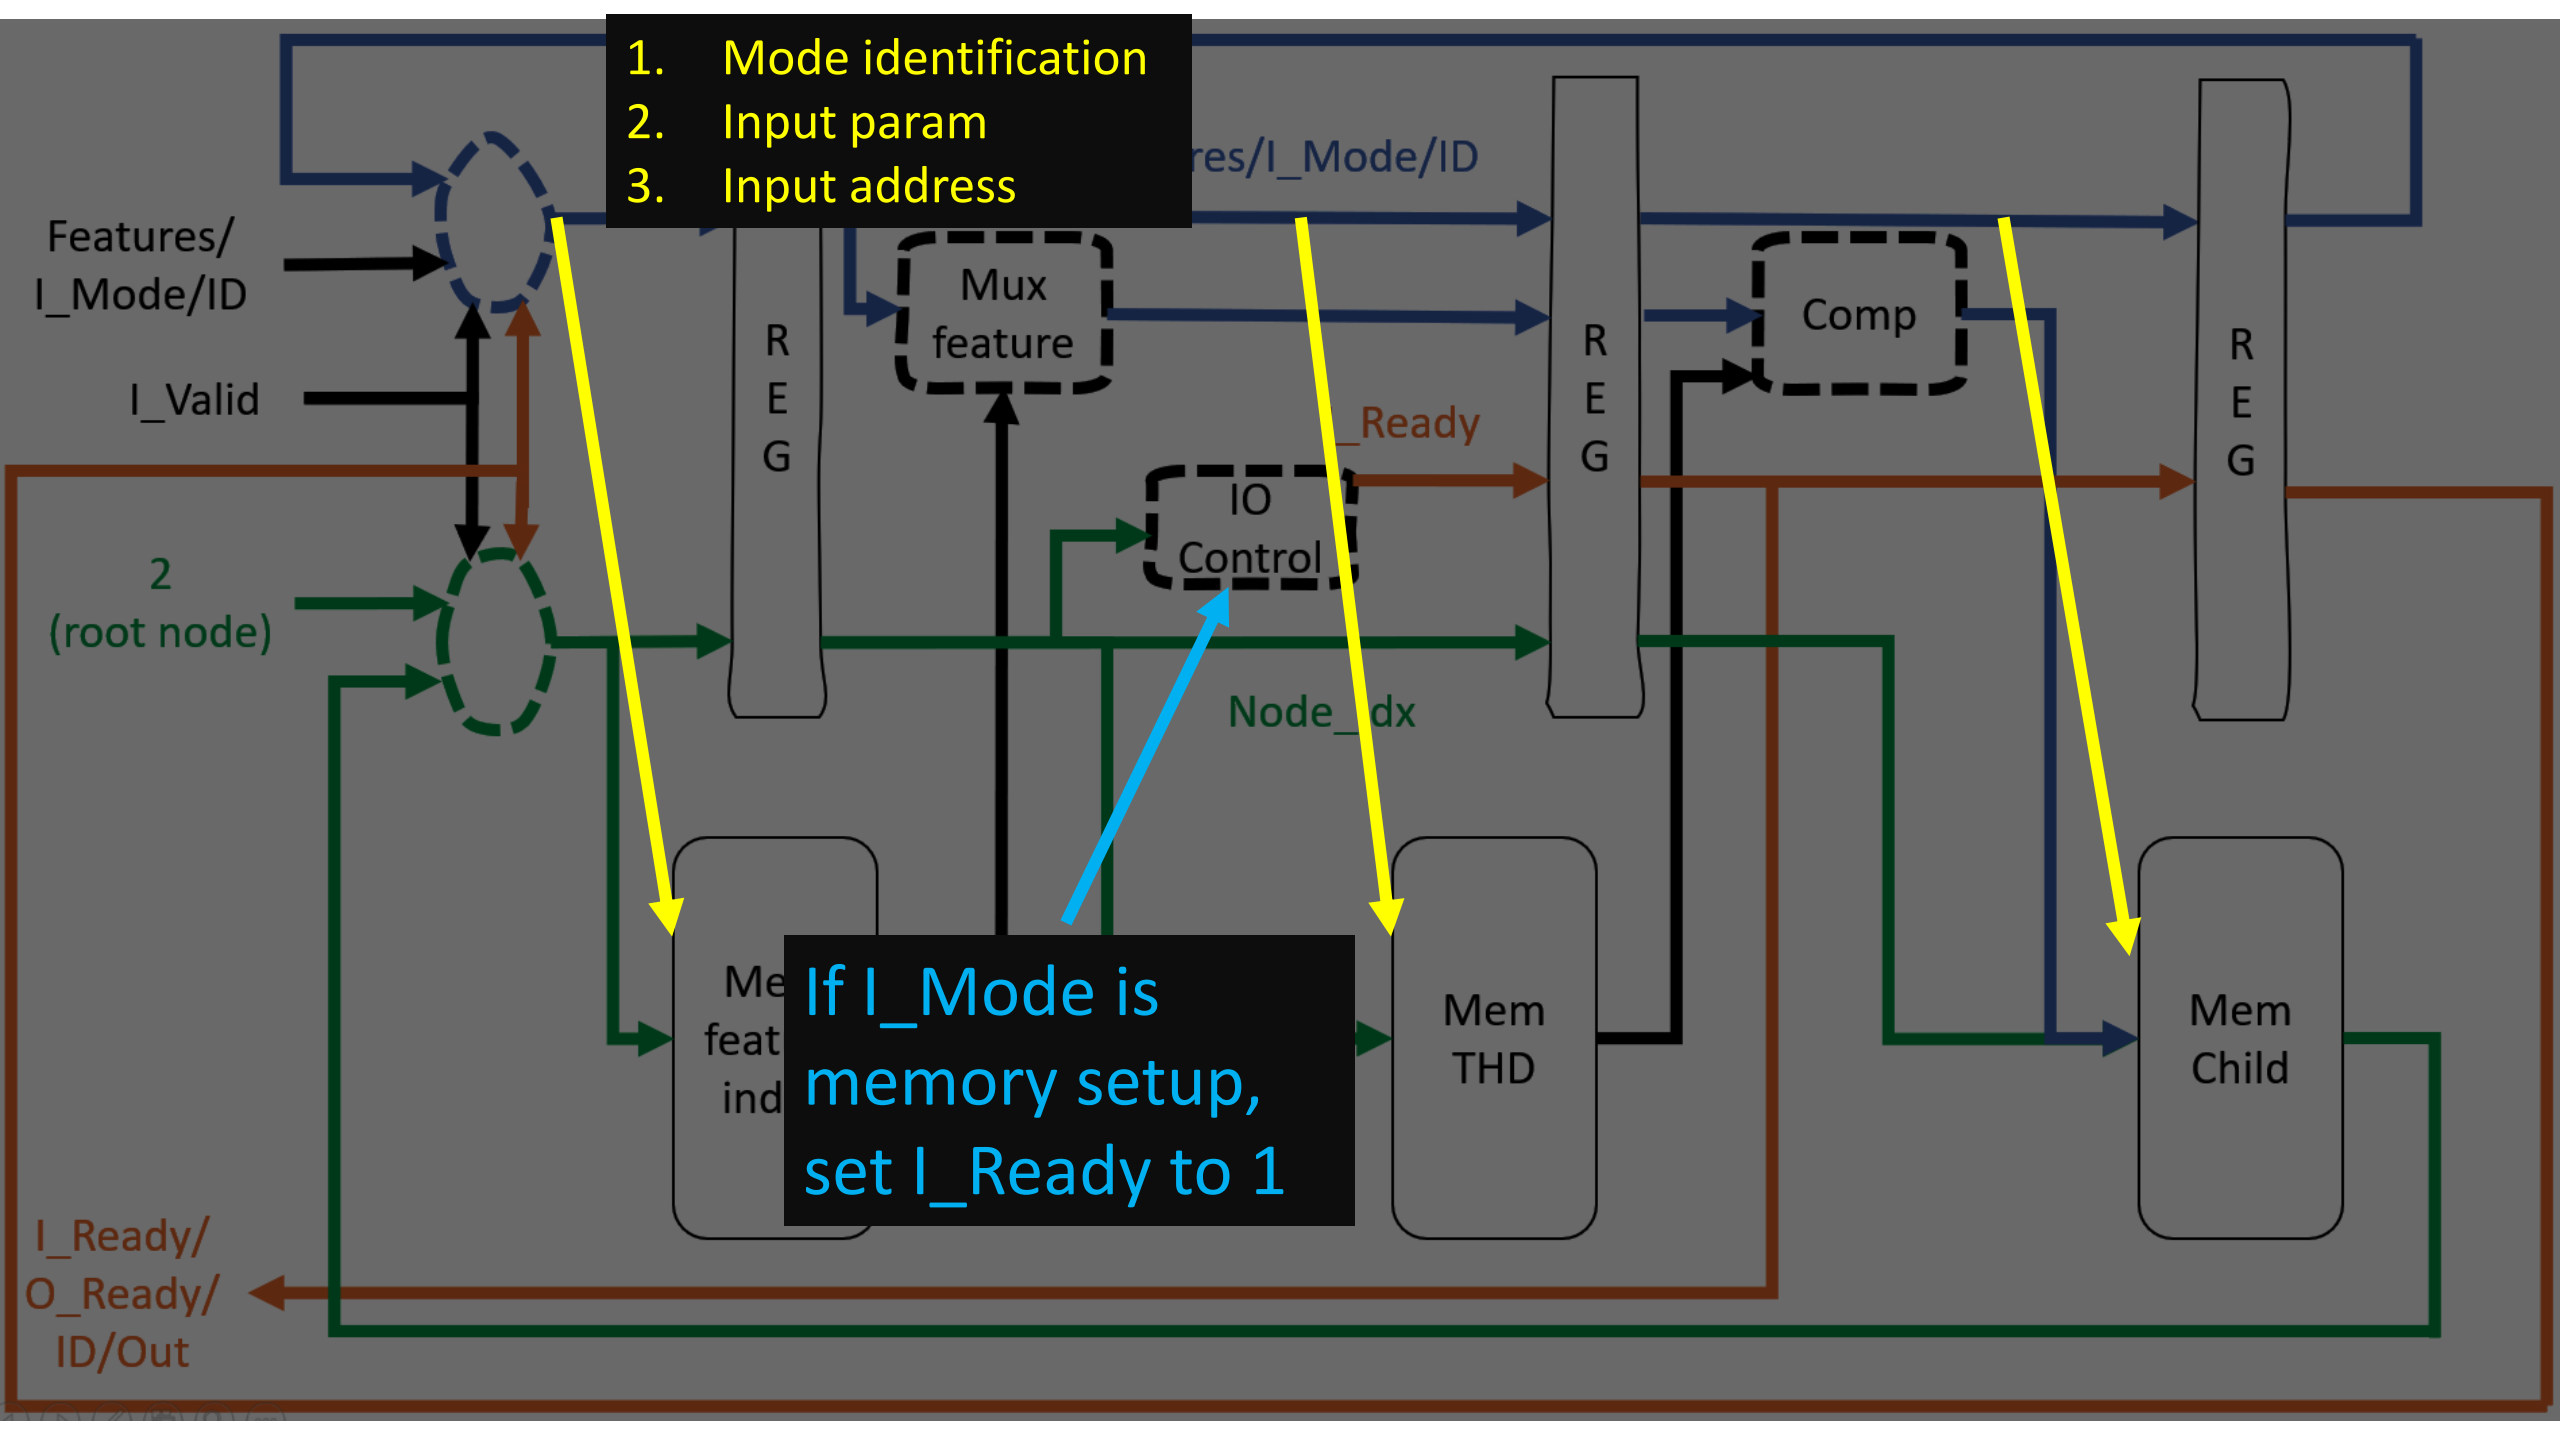






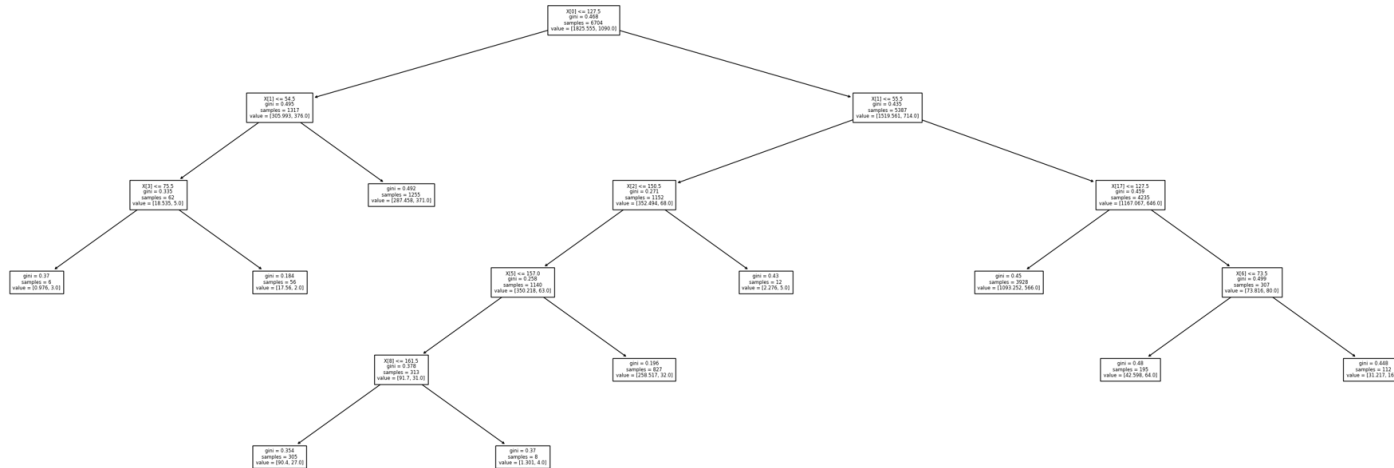
# Architecture – A deeper look

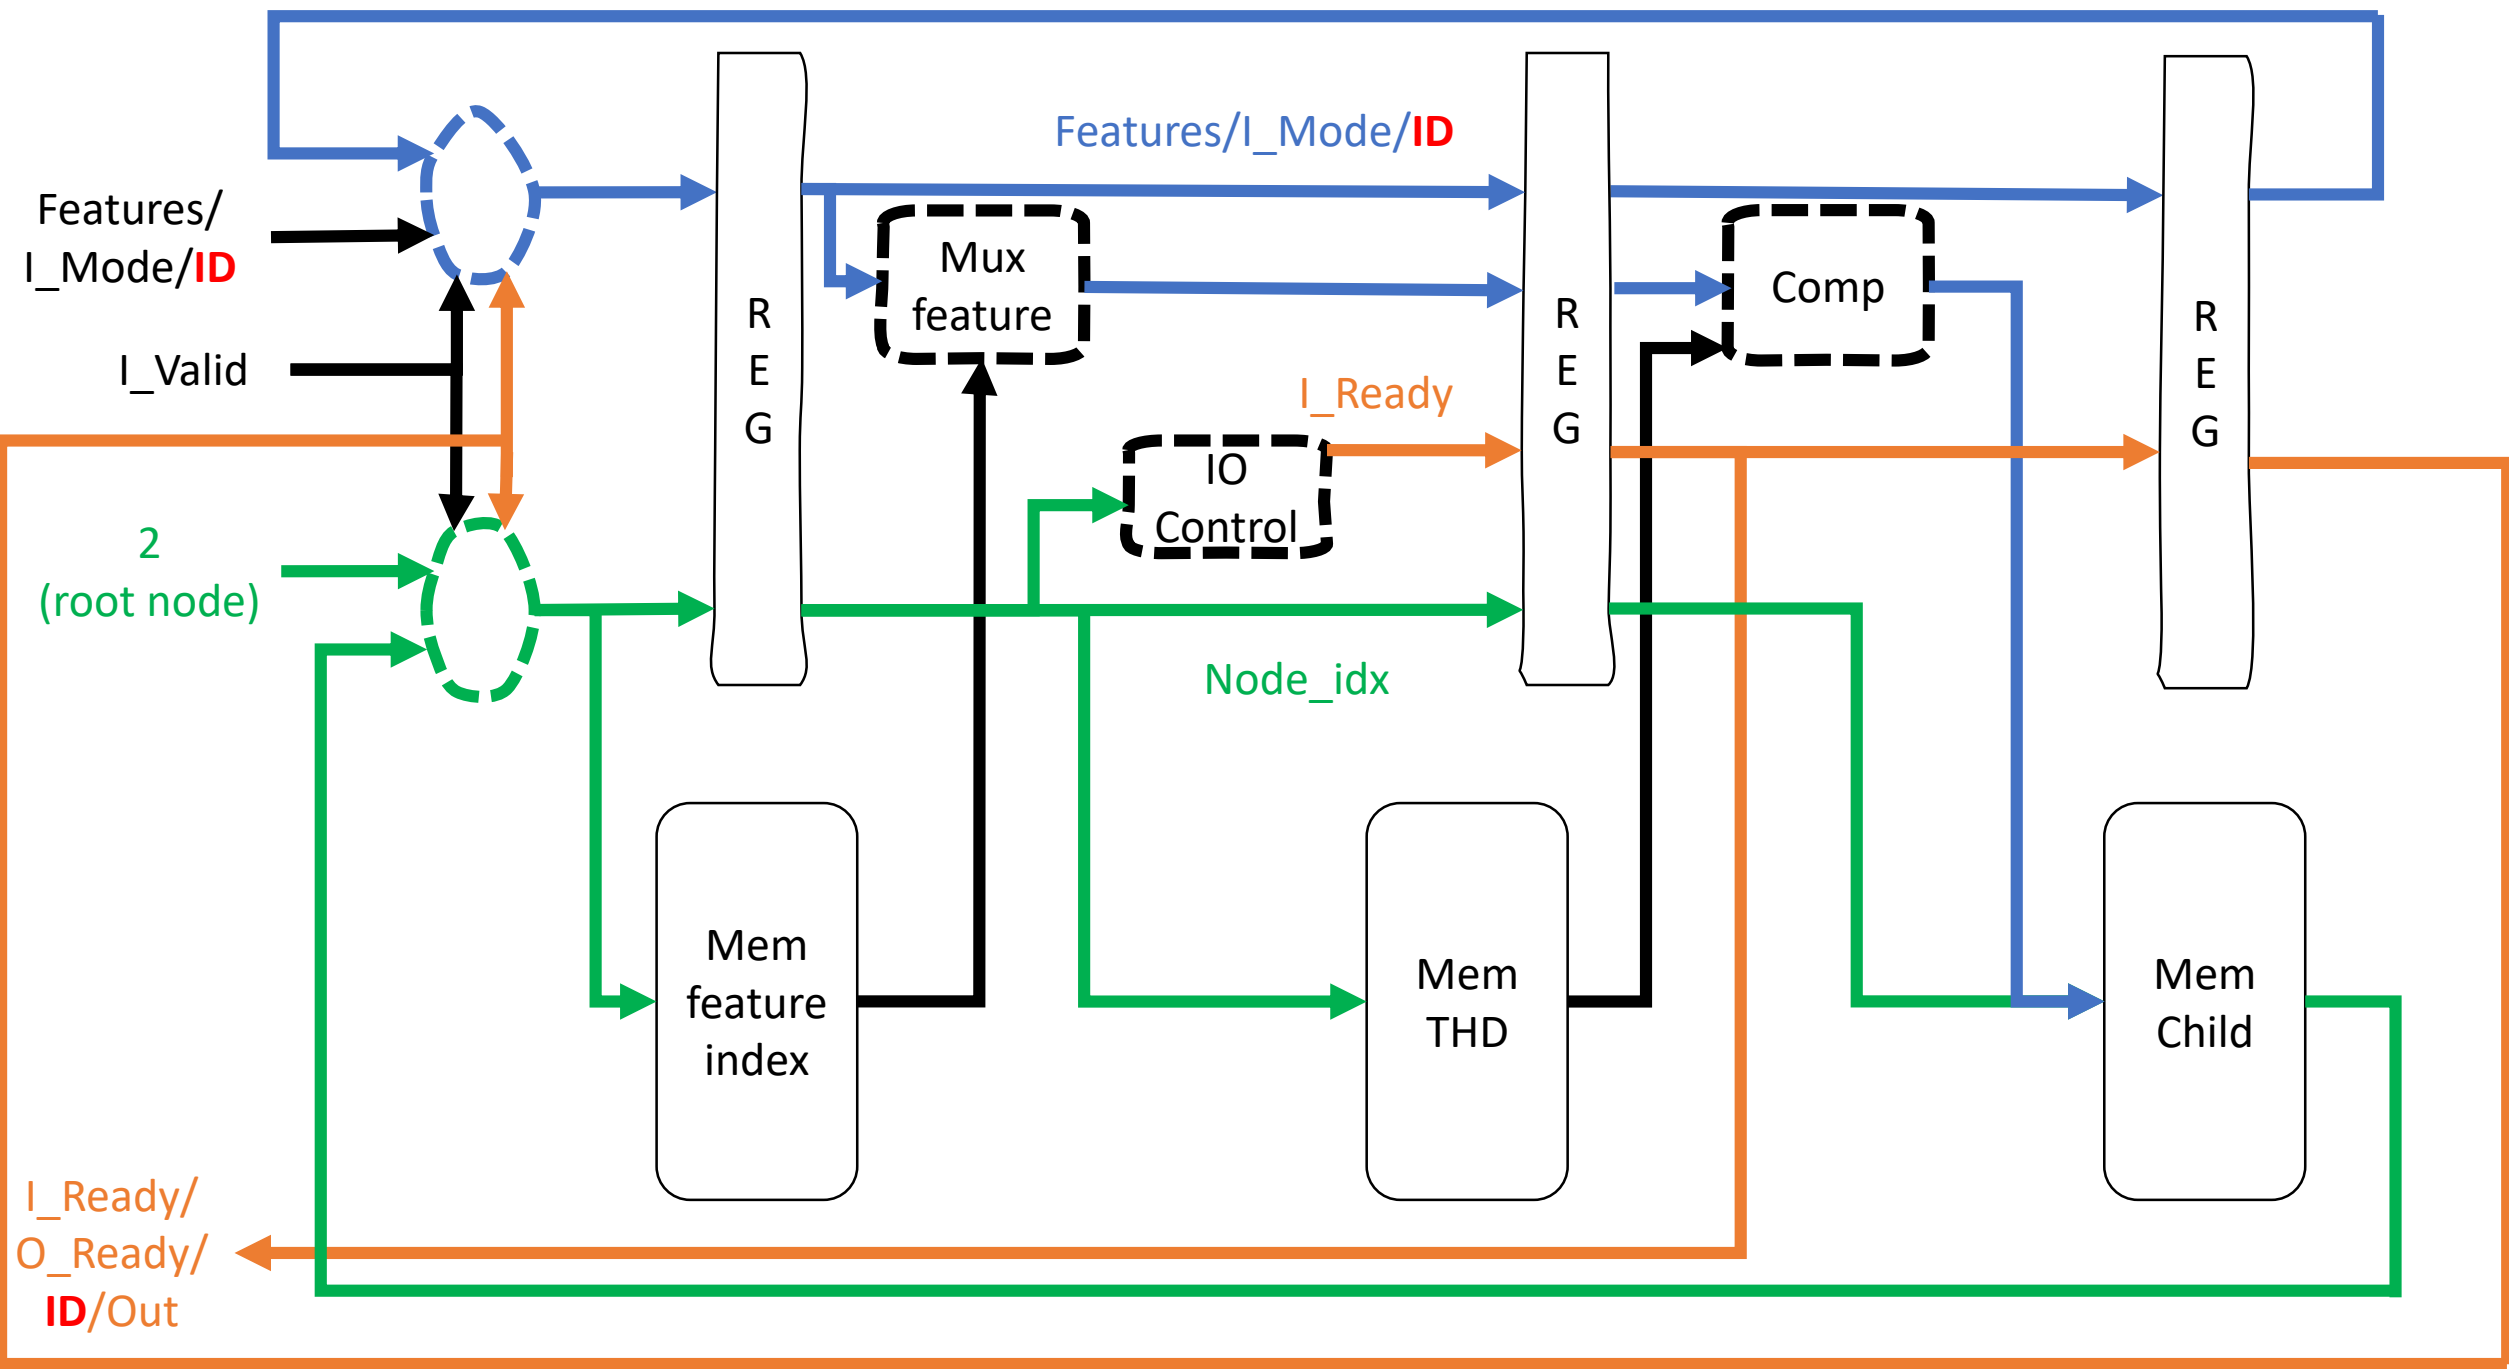
- Memory Initialization
  - Pass the params and addresses one by one
    - Re-using existing input ports
  - How to identify the correct instruction?
    - Add a new input port, I\_Mode
    - 00: Initialize feature index memory
    - 01: Initialize threshold memory
    - 10: Initialize child node memory
    - 11: Computation
  - Allows parameter modification at run time!



# Architecture – Final thoughts

- Processing latency varies between data
  - IO does not follow first-in-first-out
  - Naïve solution:
    - Track the data ID in each stage





# Synthesis

- Using Synopsys Design Vision
  - Process node: TSMC 0.13um
  - Clock Cycle: 3.6ns
  - Memory: 86% of the cell area

```
Number of ports:          92
Number of nets:          675
Number of cells:         556
Number of combinational cells: 294
Number of sequential cells: 259
Number of macros/black boxes: 3
Number of buf/inv:       44
Number of references:     58

Combinational area:      2946.686373
Buf/Inv area:           417.560399
Noncombinational area:   8352.905128
Macro/Black Box area:    72567.695312
Net interconnect area:   74480.387431

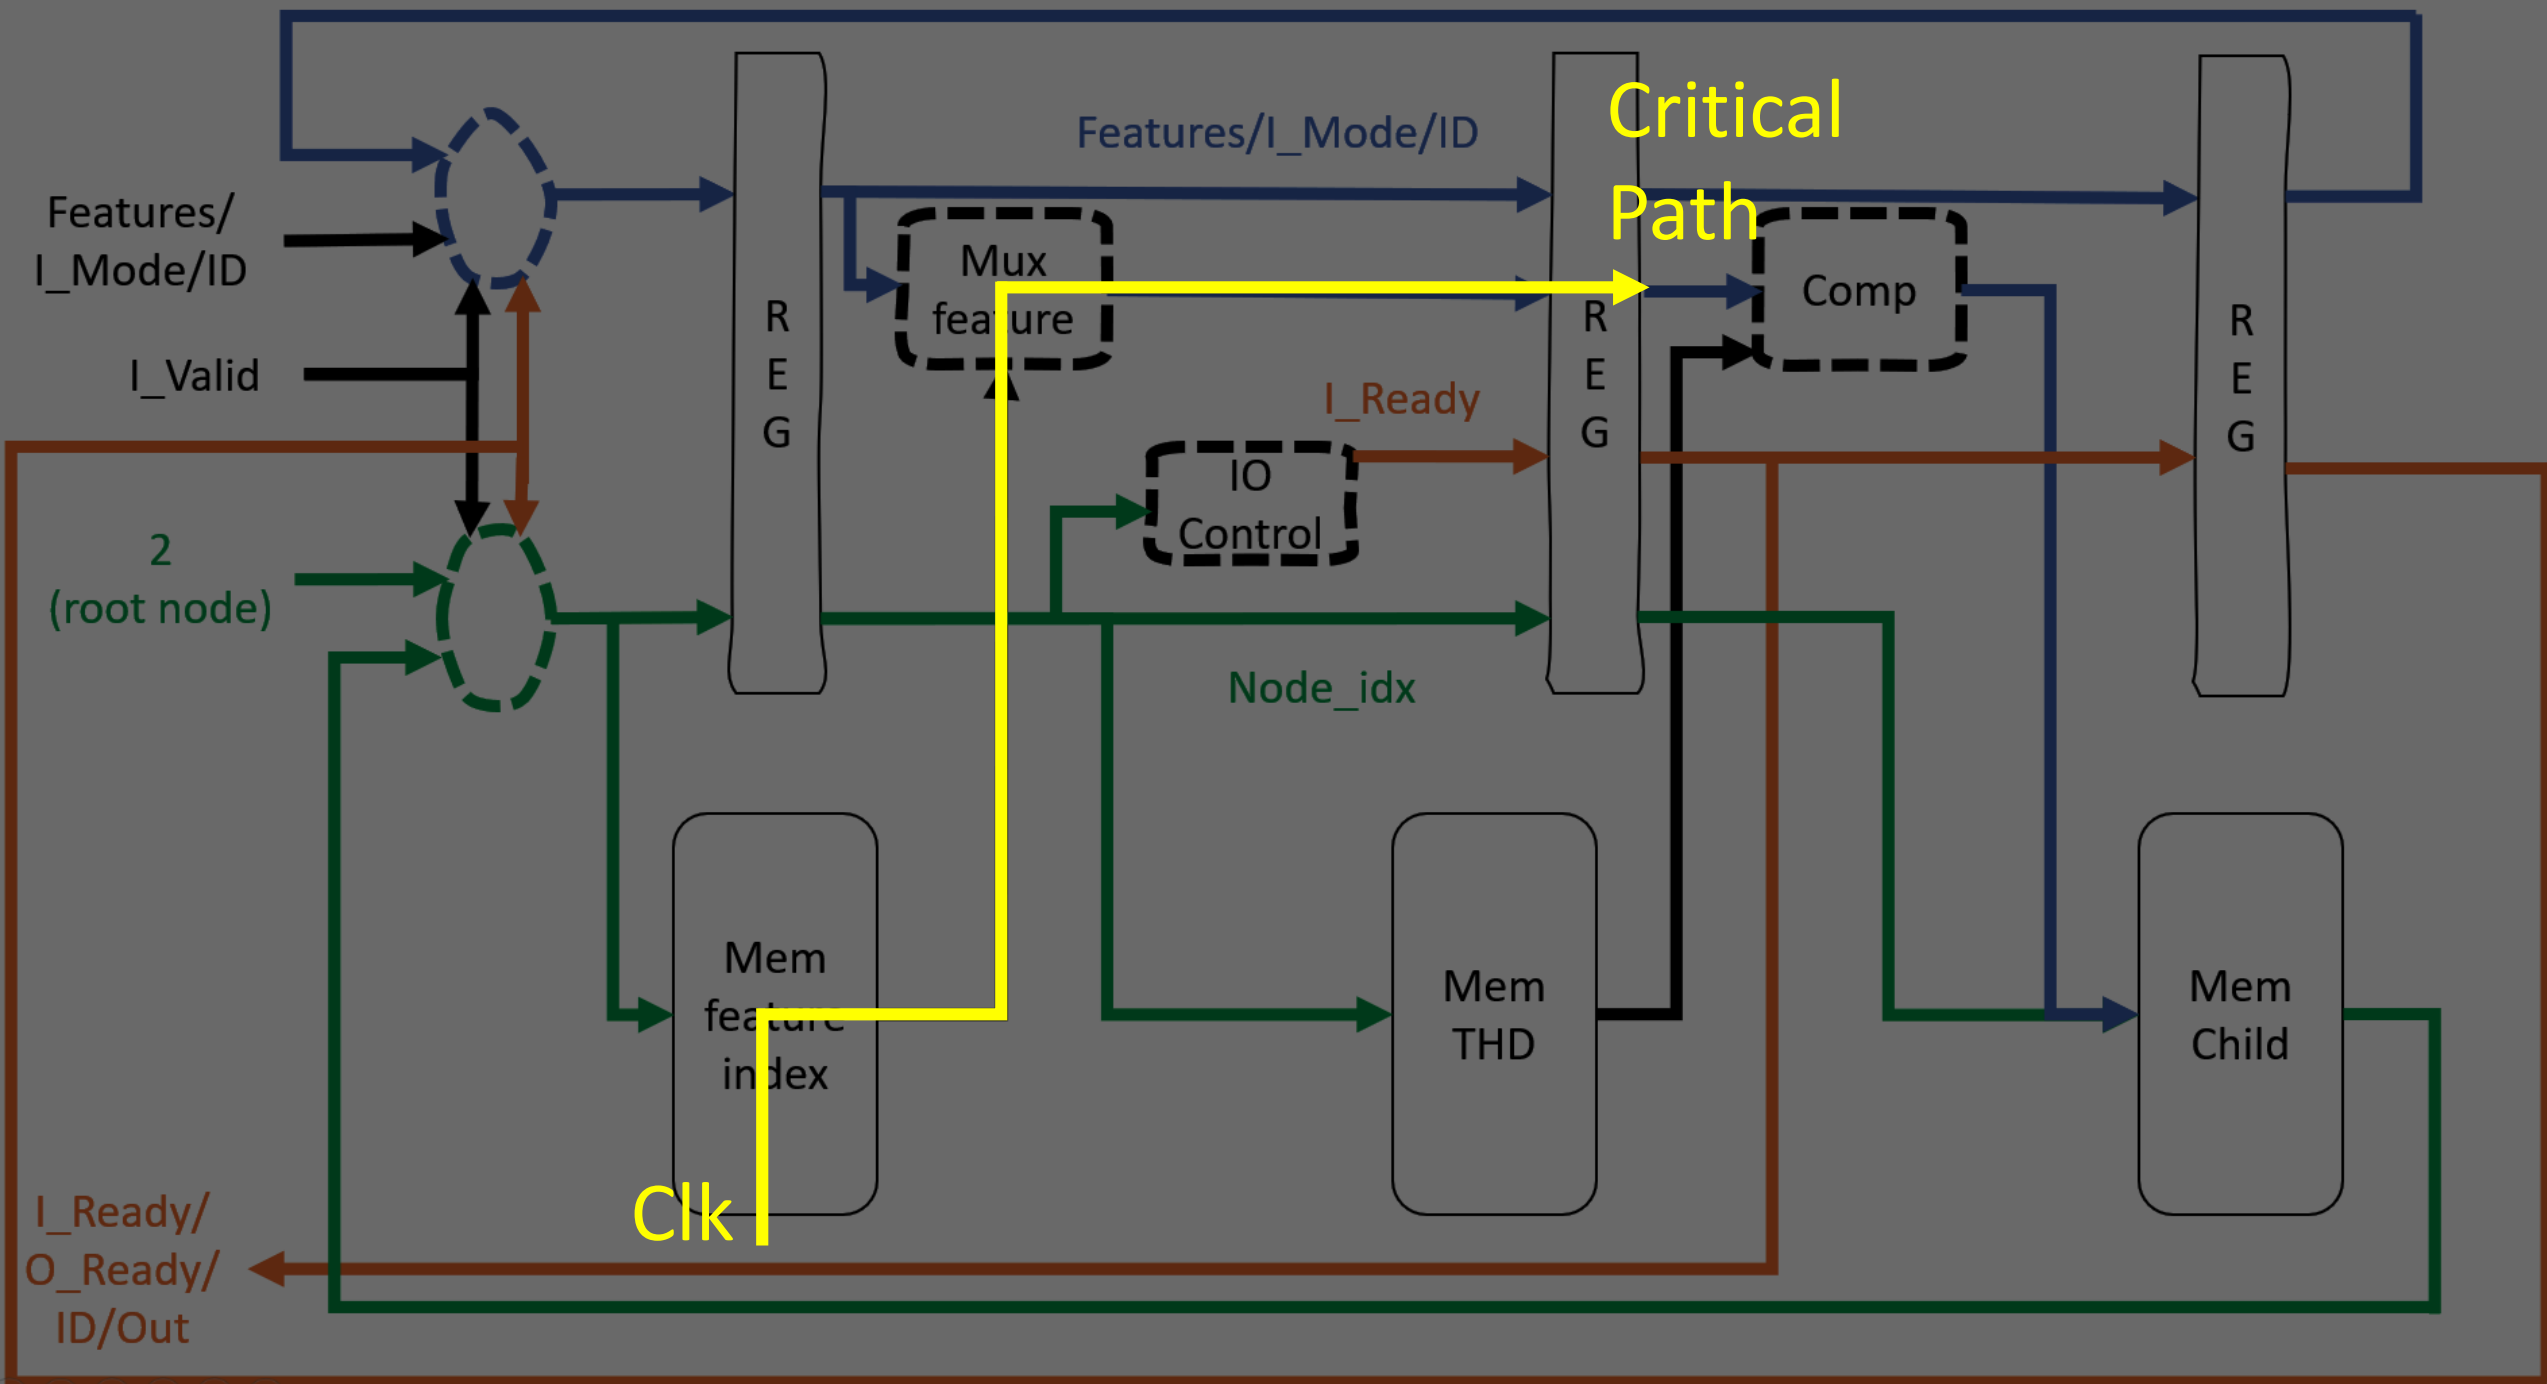
Total cell area:         83867.286814
Total area:              158347.674265
```

```
Startpoint: fea_idx_sram
             (rising edge-triggered flip-flop clocked by clk)
Endpoint: selected_feature_s2_r_reg_0_
          (rising edge-triggered flip-flop clocked by clk)
```

```
Path Group: clk
Path Type: max
```


Des/Clust/Port	Wire Load Model	Library
DEC	tsmc13_wl10	slow

Point	Incr	Path
clock clk (rise edge)	0.00	0.00
clock network delay (ideal)	0.50	0.50
fea_idx_sram/CLK (sram_256x8)	0.00	0.50 r
fea_idx_sram/Q[1] (sram_256x8)	2.06	2.56 f
U400/Y (NOR2BX4)	0.17	2.73 f
U401/Y (BUFX12)	0.19	2.92 f
U414/Y (AOI22X1)	0.29	3.21 r
U415/Y (OAI211X1)	0.23	3.44 f
U416/Y (NAND2X1)	0.22	3.66 r
U371/Y (OAI21X2)	0.12	3.78 f
selected_feature_s2_r_reg_0_/D (DFFRX1)	0.00	3.78 f
data arrival time		3.78




# Place and Route

- Using Cadence Innovus
  - Utilization Rate: 0.65
  - Memory dominating layout
    - **Floorplan, memory placement, power plan** are critical to avoid violations!
  - Pass P&R testbench simulation at clock cycle 3.6 ns





# Place and Route



# Accelerator Performance

- The evaluation metric is the inference time of all 2874 testing data.
- The baseline inference time of scikit-learn decision tree, using CPU, is 1.58 ms
  - Intel(R) Core(TM) i7-7700 CPU @ 3.60GHz

```
Total inference time using cpu: 0.0015842914581298828
```

- The presented decision tree accelerator only requires 42.5 us!
  - Including memory setup
  - ~37 times faster!

```
PASS  
Simulation complete via $finish(1) at time 42508800 PS + 0  
./tb_v3.v:226 $finish;
```

# References

1. <https://github.com/Shayan-Asgari/ClassificationTrees>
2. <https://stackoverflow.com/questions/51397109/prune-unnecessary-leaves-in-sklearn-decisiontreeclassifier>
3. <https://stackoverflow.com/questions/56334210/how-to-extract-sklearn-decision-tree-rules-to-pandas-boolean-conditions>
4. Course material of COMPUTER-AIDED VLSI SYSTEM DESIGN,  
National Taiwan University

# About me

- Bo-Fan, Chen
- Nationality: Taiwan
- Language
  - Mandarin
  - English
- Education
  - National Taiwan University
    - Electric Engineering, graduated in 2019.07
    - Graduate Institute of Electronics Engineering, Integrated Circuits & Systems Group, Laboratory for Data Processing Systems, expect to graduate in 2022.09
  - Aarhus University, Denmark
    - Exchange student, from 2021.09 to 2022.07
    - Software Engineering, Electric Engineering
- Academic Work
  - CF-NET: Complementary Fusion Network for Rotation Invariant Point Cloud Completion, ICASSP 2022
- Skills
  - Digital Circuit Design, Verilog, Synopsys Design Vision, Cadence Innovus
  - Machine Learning, Python, Numpy, Pytorch, Tensorflow
- Objective
  - Looking for a position as digital circuit RTL engineer. Preferred locations are: United States, Europe, Taipei.



# Contact

- Email
  - [sddslover@gmail.com](mailto:sddslover@gmail.com)
- LinkedIn
  - [linkedin.com/in/bo-fan-chen](https://www.linkedin.com/in/bo-fan-chen)